

Important Links:

<https://aws.amazon.com/certification/certified-cloud-practitioner/>

<https://explore.skillbuilder.aws/lms/>

Tips

1. Utilize 50% discount of exam cost
2. Study topics in order
3. Don't over cross the topics
4. Better make you notes

Which benefit of Cloud Computing allows AWS to offer lower pay-as-you-go prices as usage from hundreds of thousands of customers is aggregated **in the cloud**?

- Massive economies of scale**

(Correct)

AWS Cloud History

2002:
Internally
launched

2004:
Launched publicly
with SQS

2007:
Launched in
Europe



2003:
Amazon infrastructure is
one of their core strength.
Idea to market

2006:
Re-launched
publicly with
SQS, S3 & EC2



WHAT IS AN AVAILABILITY ZONE?

- An AWS availability zone is a logical data center that is located in a certain region.
- There are two or more availability zones in every AWS region.
- In order to ensure availability and reduce the likelihood of having two zones going down at the same time, each zone has a redundant and separate power and networking.
- *Note: A data center consists of bunch of servers*
- *Note: A common misconception is that a single zone equals a single data center. In fact, each zone is backed by one or more physical data centers.*



Good job!

AWS Regions consist of multiple, isolated, and physically separate Availability Zones within a geographic area.

Question 8:

AWS Regions are composed of?

Two or more Edge Locations

One or more discrete data centers

Two or more Availability Zones



REGION

AVAILABILITY
ZONE A

AVAILABILITY
ZONE B

AVAILABILITY
ZONE C

Which of the following S3 storage classes takes the most time to retrieve data (also known as first byte latency)?

A photo sharing web application wants to store thumbnails of user-uploaded images on Amazon S3. The thumbnails are rarely used but need to be immediately accessible from the web application. The thumbnails can be regenerated easily if they are lost. Which is the most cost-effective way to store these thumbnails on S3?



Use S3 One-Zone Infrequent Access (One-Zone IA) to store the thumbnails

(Correct)

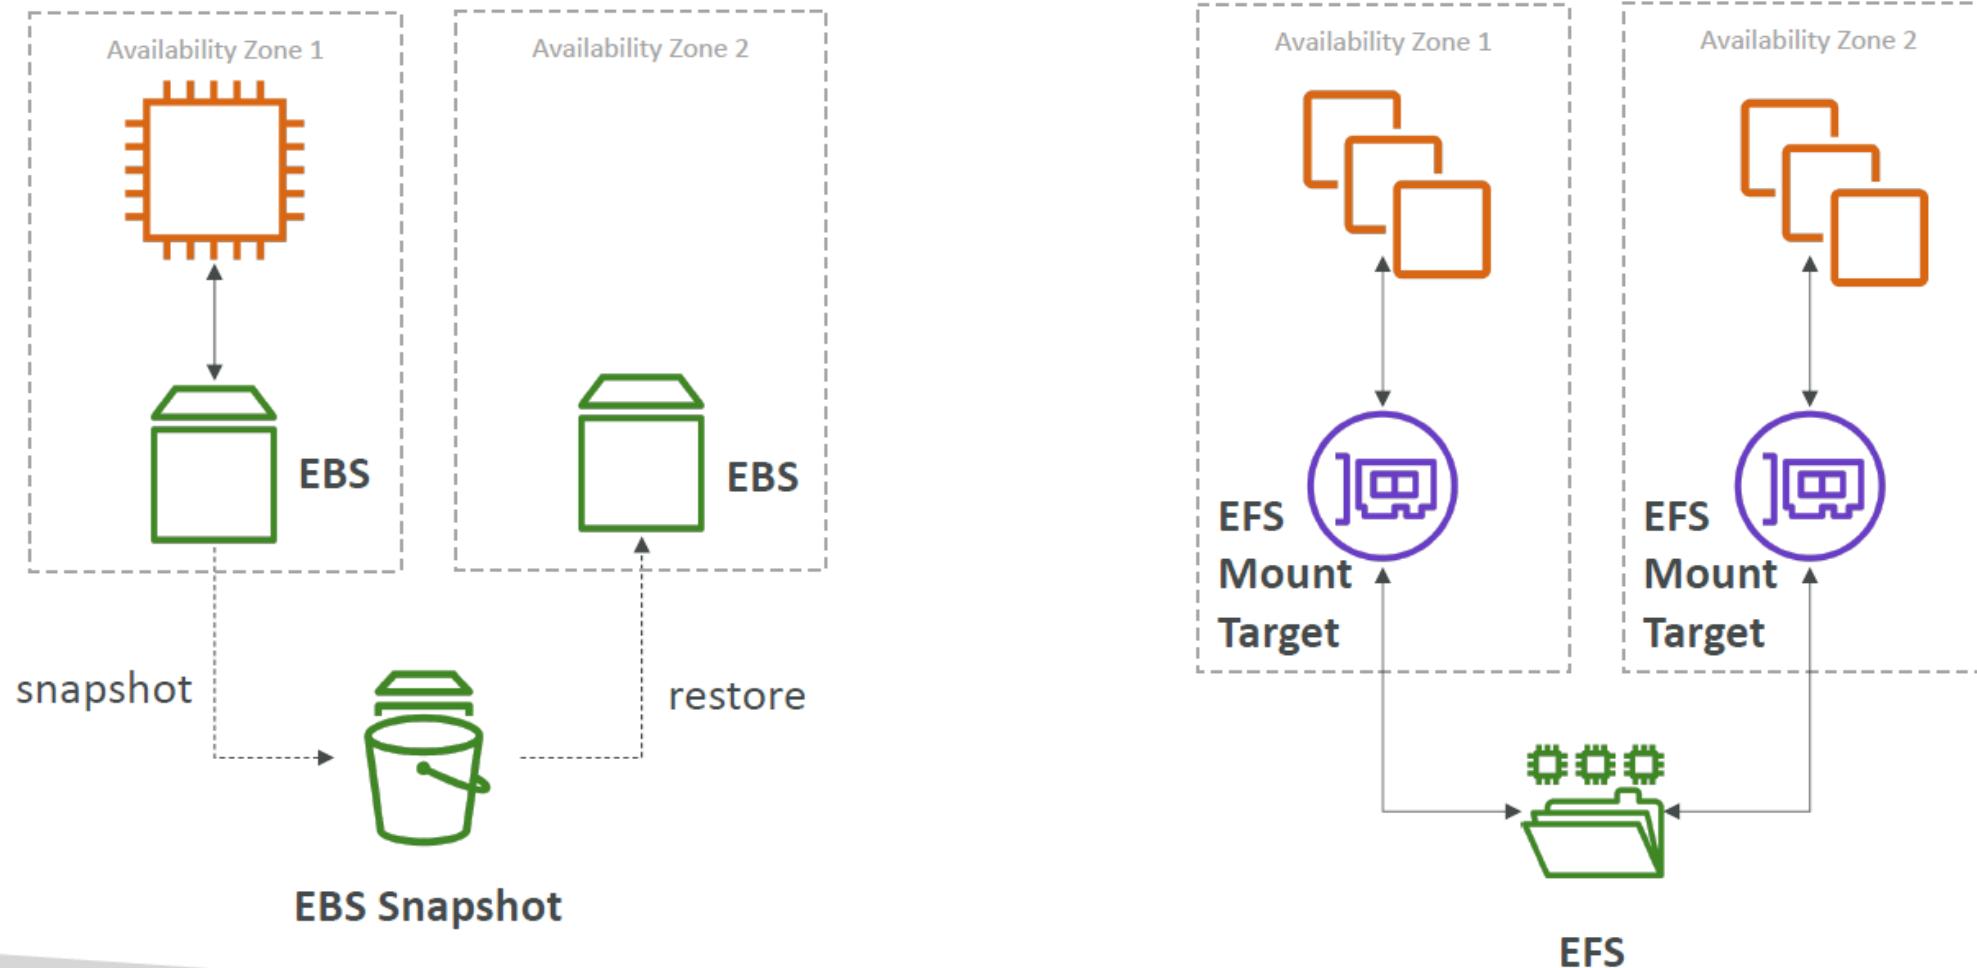
Performance across the S3 Storage Classes

Which S3 storage class offers the lowest availability?

Which of the following S3 storage classes has NO constraint of a minimum **storage duration charge for objects?** **S3 STANDARD**

	S3 Standard	S3 Intelligent-Tiering*	S3 Standard-IA	S3 One Zone-IA†	S3 Glacier Instant Retrieval	S3 Glacier Flexible Retrieval	S3 Glacier Deep Archive
Designed for durability	99.999999999% (11 9's)	99.999999999% (11 9's)	99.999999999% (11 9's)	99.999999999% (11 9's)	99.999999999% (11 9's)	99.999999999% (11 9's)	99.999999999% (11 9's)
Designed for availability	99.99%	99.9%	99.9%	99.5%	99.9%	99.99%	99.99%
Availability SLA	99.9%	99%	99%	99%	99%	99.%	99.9%
Availability Zones	≥3	≥3	≥3	1	≥3	≥3	≥3
Minimum capacity charge per object	N/A	N/A	128 KB	128 KB	128 KB	40 KB	40 KB
Minimum storage duration charge	N/A	N/A	30 days	30 days	90 days	90 days	180 days
Retrieval charge	N/A	N/A	per GB retrieved	per GB retrieved	per GB retrieved	per GB retrieved	per GB retrieved
First byte latency	milliseconds	milliseconds	milliseconds	milliseconds	milliseconds	minutes or hours	hours
Storage type	Object	Object	Object	Object	Object	Object	Object
Lifecycle transitions	Yes	Yes	https://www.youtube.com/@atozaboutdata/	Yes	Yes	Yes	Yes

EBS vs EFS



S3 is a key value based object storage service

S3 stores data in a flat non-hierarchical structure

Amazon Simple Storage Service (Amazon S3) is an object storage service that offers industry-leading scalability, data availability, security, and performance. S3 stores data in a flat non-hierarchical structure. All objects are stored in S3 buckets and can be organized with shared names called prefixes. You can also append up to 10 key-value pairs called S3 object tags to each object, which can be created, updated, and deleted throughout an object's lifecycle.

Data transferred in from the internet

Data transferred out to an Amazon Elastic Compute Cloud (Amazon EC2) instance, when the instance is in the same AWS Region as the S3 bucket

There are four cost components to consider for S3 pricing – **storage pricing; request and data retrieval pricing; data transfer and transfer acceleration pricing**; and data management features pricing. Under "Data Transfer", You pay for all bandwidth into and out of Amazon S3, except for the following:

(1) Data transferred in from the internet, (2) Data transferred out to an Amazon Elastic Compute Cloud (Amazon EC2) instance, when the instance is in the same AWS Region as the S3 bucket, (3) Data transferred out to Amazon CloudFront (CloudFront).

Each AWS Region consists of two or more Availability Zones

Each Availability Zone (AZ) consists of one or more discrete data centers

AWS has the concept of a Region, which is a physical location around the world where AWS clusters data centers. Each AWS Region consists of multiple (two or more), isolated, and physically separate AZ's within a geographic area. Each AZ has independent power, cooling, and physical security and is connected via redundant, ultra-low-latency networks.

An Availability Zone (AZ) is one or more discrete data centers with redundant power, networking, and connectivity in an AWS Region. All AZ's in an AWS Region are interconnected with high-bandwidth, low-latency networking, over fully redundant, dedicated metro fiber providing high-throughput, low-latency networking between AZ's.

With AWS Config, you can do the following: 1. Evaluate your AWS resource configurations for desired settings. 2. Get a snapshot of the current configurations of the supported resources that are associated with your AWS account. 3. Retrieve configurations of one or more resources that exist in your account. 4. Retrieve historical configurations of one or more resources. 5. Receive a notification whenever a resource is created, modified, or deleted. 6. View relationships between resources. For example, **you might want to find all resources that use a particular security group**.

AWS Shield, WAF, and CloudFront with Route 53 help mitigate DDoS attacks:

Which of the following statements are CORRECT regarding the AWS VPC service?
(Select two)

A Security Group can have allow rules only

A NAT Gateway is managed by AWS

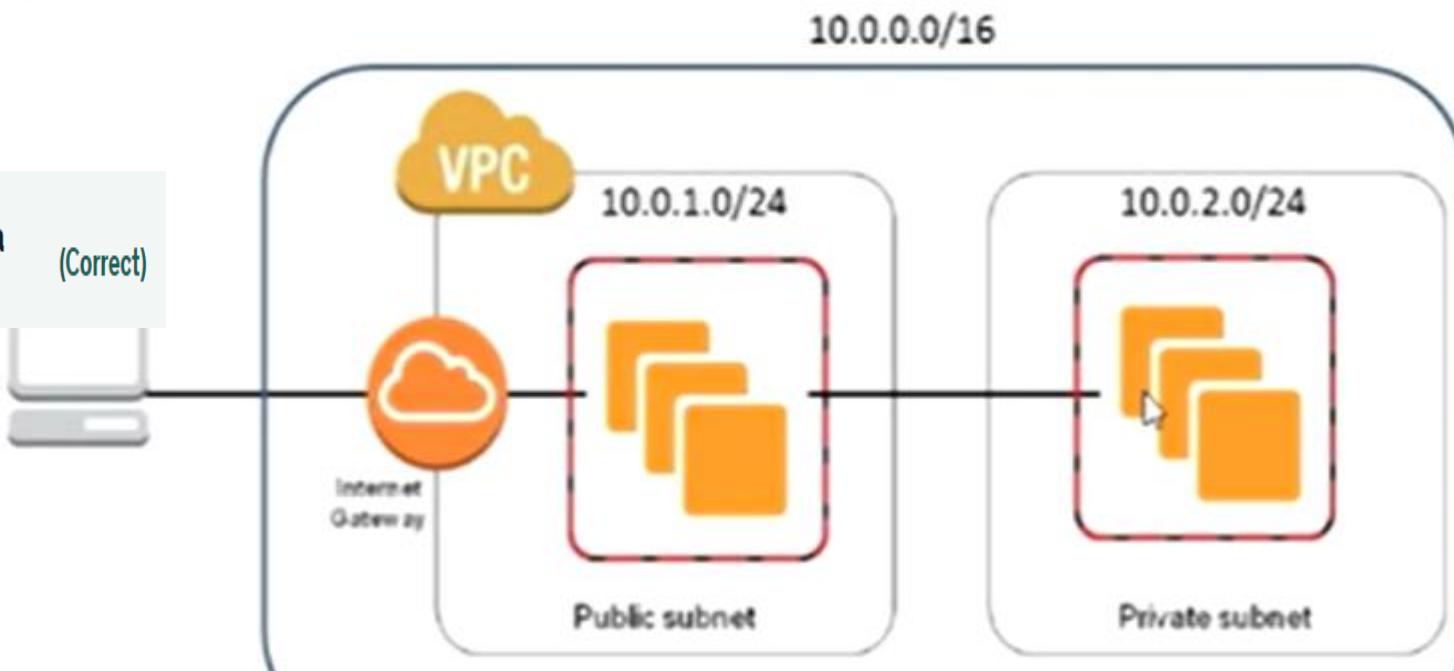
A Network Access Control List (NACL) is an optional layer of security for your VPC that acts as a firewall for controlling traffic in and out of one or more subnets (i.e. it works at subnet level). A network ACL has separate inbound and outbound rules, and each rule can either allow or deny traffic.

AVPC spans all of the Availability Zones in the Region whereas a subnet spans only one Availability Zone in the Region

Security group basics

The following are the basic characteristics of security groups for your VPC:

- There are quotas on the number of security groups that you can create per VPC, the number of rules that you can add to each security group, and the number of security groups that you can associate with a network interface. For more information, see [Amazon VPC quotas](#).
- You can specify **allow rules**, but not **deny rules**.
- You can specify **separate rules for inbound and outbound traffic**.
- When you create a security group, it has no inbound rules. Therefore, no inbound traffic originating from another host to your instance is allowed until you add inbound rules to the security group.
- By default, a security group includes an outbound rule that allows all outbound traffic. You can remove the rule and add outbound rules that allow specific outbound traffic only. If your security group has no outbound rules, no outbound traffic originating from your instance is allowed.
- Security groups are stateful — if you send a request from your instance, the response traffic for that request is allowed to flow in regardless of inbound security group rules. Responses to allowed inbound traffic are allowed to flow out, regardless of outbound rules.



Which of the following AWS services support VPC Endpoint Gateway for a private connection from a VPC? (Select two)

A VPC endpoint enables **you to privately connect your VPC to supported AWS services and VPC endpoint services powered by AWS PrivateLink without requiring an internet gateway**, NAT device, VPN connection, or AWS Direct Connect connection. Instances in your VPC do not require public IP addresses to communicate with resources in the service. Traffic between your VPC and the other service does not leave the Amazon network.

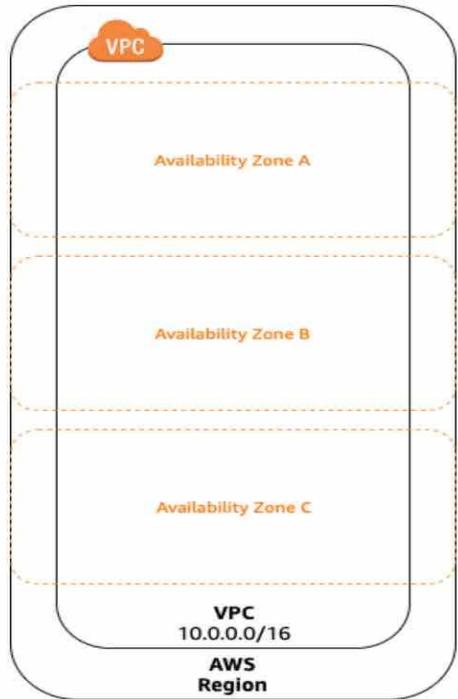
There are two types of VPC endpoints: interface endpoints and gateway endpoints.

An interface endpoint is an elastic network interface with a private IP address from the IP address range of your subnet that serves as an entry point for traffic destined to a supported service. Interface endpoints are powered by AWS PrivateLink, a technology that enables you to privately access services by using private IP addresses.

A gateway endpoint is a gateway that you specify as a target for a route in your route table for traffic destined to a supported AWS service. Just remember that **only S3 and DynamoDB support VPC Endpoint Gateway**. All other services that support VPC Endpoints use a VPC Endpoint Interface.

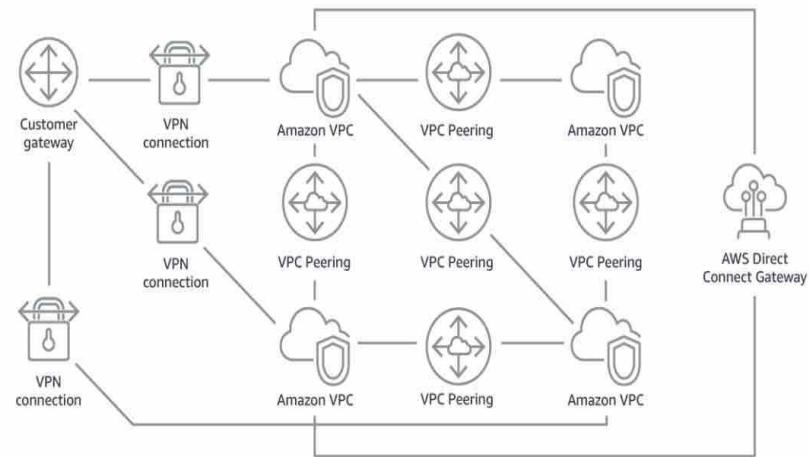
Which of the following AWS services can be used to connect a company's on-premises environment to a VPC without using the public internet?

The following diagram shows a new VPC with an IPv4 CIDR block.

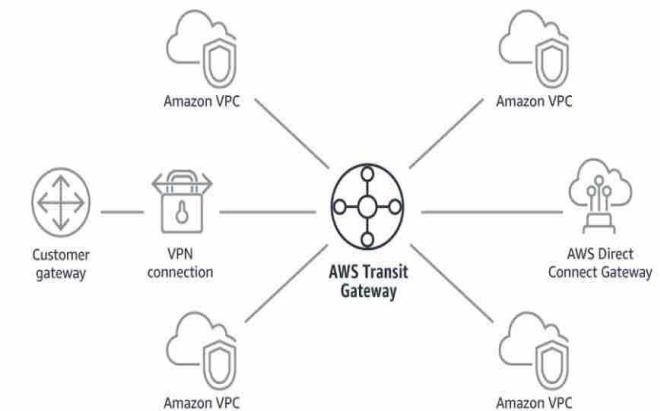


0 A VPC spans all of the Availability Zones in the Region whereas a subnet spans only one Availability Zone in the Region (Correct)

Without AWS Transit Gateway



With AWS Transit Gateway



The main route table has the following routes.

Destination	Target
10.0.0.0/16	local

Complexity increases with scale. You must maintain routing tables within each VPC and connect to each onsite location using separate network gateways.

Your network is streamlined and scalable. AWS Transit Gateway routes all traffic to and from each VPC or VPN, and you have one place to manage and monitor it all.

VPC Peering - A VPC peering connection is a networking connection between two VPCs that enables you to route traffic between them privately.

A VPC spans all of the Availability Zones in the Region. After creating a VPC, you can add one or more subnets in each Availability Zone. You can optionally add subnets in a Local Zone, which is an AWS infrastructure deployment that places compute, storage, database, and other select services closer to your end users. A Local Zone enables your end users to run applications that require single-digit millisecond latencies. For information about the Regions that support Local Zones, see [Available Regions](#) in the *Amazon EC2 User Guide for Linux Instances*. When you create a subnet, you specify the CIDR block for the subnet, which is a subset of the VPC CIDR block. Each subnet must reside entirely within one Availability Zone and cannot span zones. Availability Zones are distinct locations that are engineered to be isolated from failures in other Availability Zones. By launching instances in separate Availability Zones, you can protect your applications from the failure of a single location. We assign a unique ID to each subnet.

Components of your Site-to-Site VPN

A Site-to-Site VPN connection offers two VPN tunnels between a virtual private gateway or a transit gateway on the AWS side, and a customer gateway on the remote (on-premises) side.

A Site-to-Site VPN connection consists of the following components. For more information about Site-to-Site VPN quotas, see [Site-to-Site VPN quotas](#).

Contents

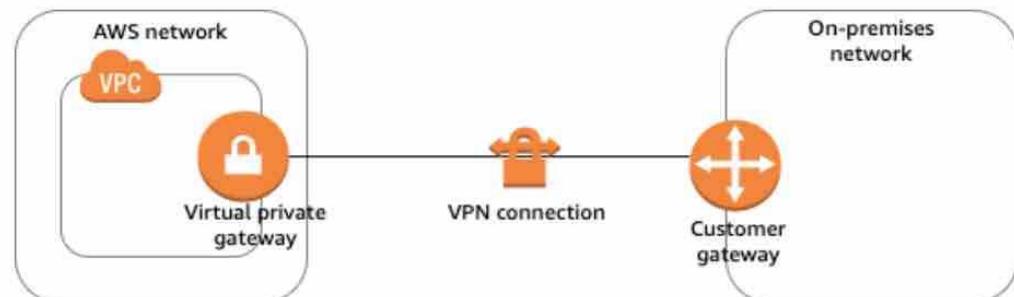
- [Virtual private gateway](#)
- [Transit gateway](#)
- [Customer gateway](#)
- [Customer gateway device](#)

Which of the following are components of an AWS Site-to-Site VPN? (Select two)

- | | |
|---|-----------|
| <input type="checkbox"/> Customer Gateway | (Correct) |
| <input type="checkbox"/> Virtual Private Gateway | (Correct) |

Virtual private gateway

A *virtual private gateway* is the VPN concentrator on the Amazon side of the Site-to-Site VPN connection. You create a virtual private gateway and attach it to the VPC from which you want to create the Site-to-Site VPN connection.



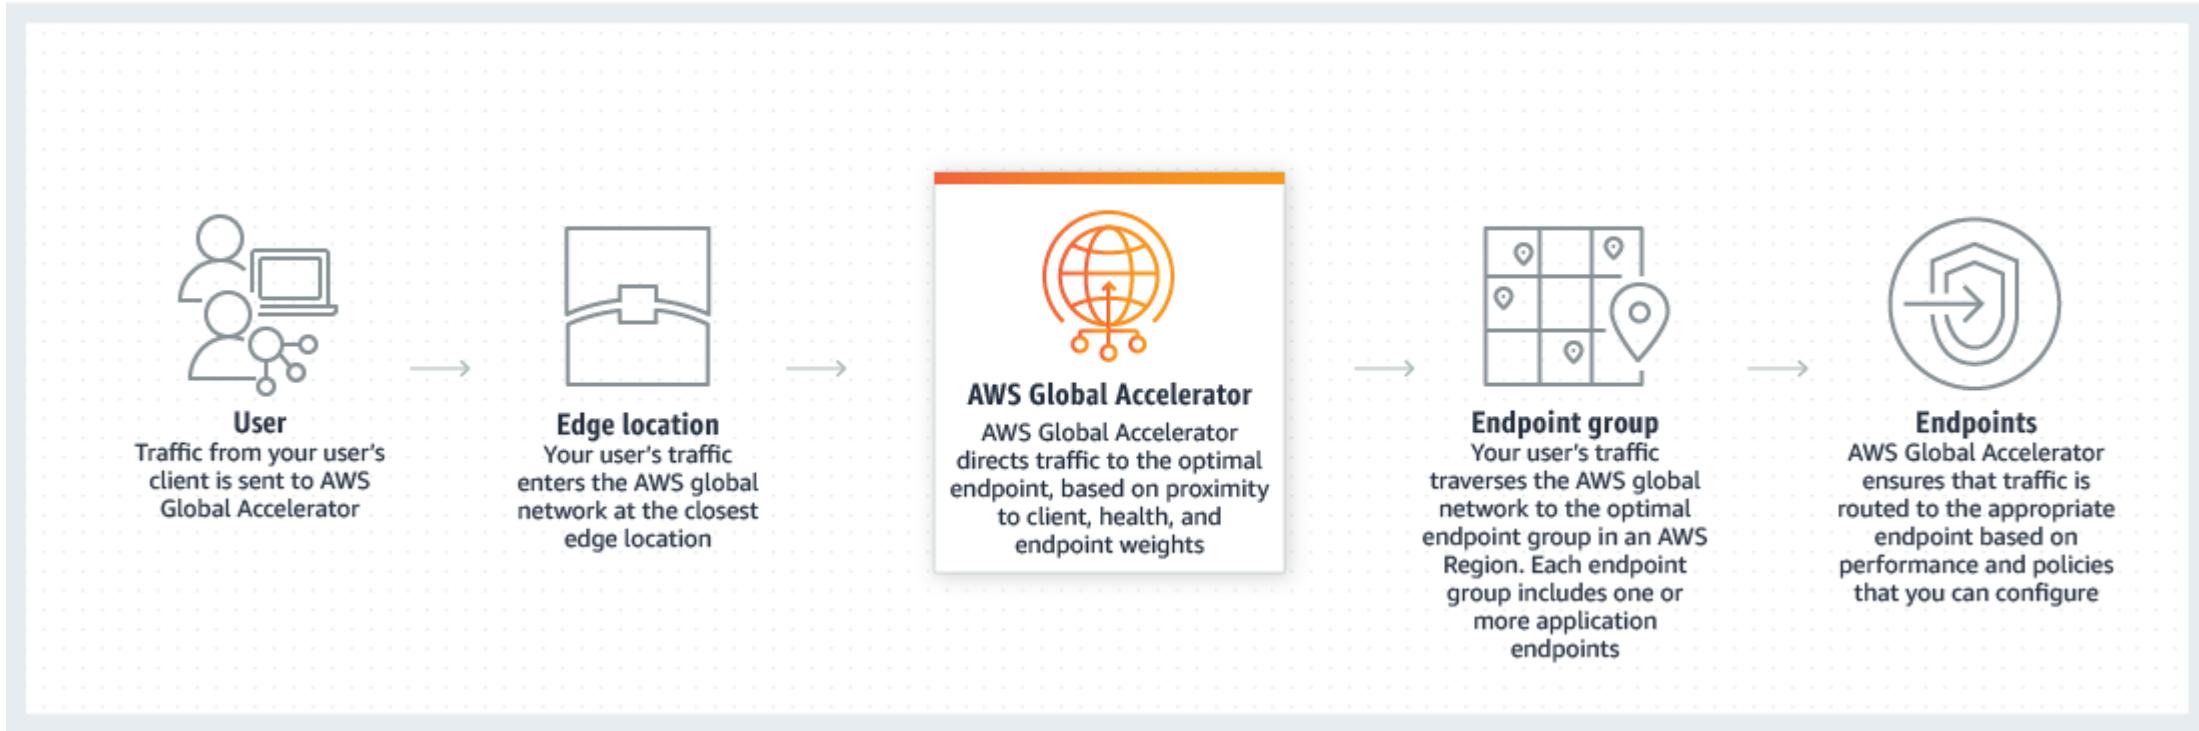
EC2 Instance Storage - Summary

- **EBS volumes:**
 - network drives attached to one EC2 instance at a time
 - Mapped to an Availability Zones
 - Can use EBS Snapshots for backups / transferring EBS volumes across AZ
- **AMI:** create ready-to-use EC2 instances with our customizations
- **EC2 Image Builder:** automatically build, test and distribute AMIs
- **EC2 Instance Store:**
 - High performance hardware disk attached to our EC2 instance
 - Lost if our instance is stopped / terminated
- **EFS:** network file system, can be attached to 100s of instances in a region
- **EFS-IA:** cost-optimized storage class for infrequent accessed files
- **FSx for Windows:** Network File System for Windows servers
- **FSx for Lustre:** High Performance Computing Linux file system

Scalability vs Elasticity (vs Agility)

- **Scalability:** ability to accommodate a larger load by making the hardware stronger (scale up), or by adding nodes (scale out)
- **Elasticity:** once a system is scalable, elasticity means that there will be some “auto-scaling” so that the system can scale based on the load. This is “cloud-friendly”: pay-per-use, match demand, optimize costs
- **Agility:** (not related to scalability - distractor) new IT resources are only a click away, which means that you reduce the time to make those resources available to your developers from weeks to just minutes.

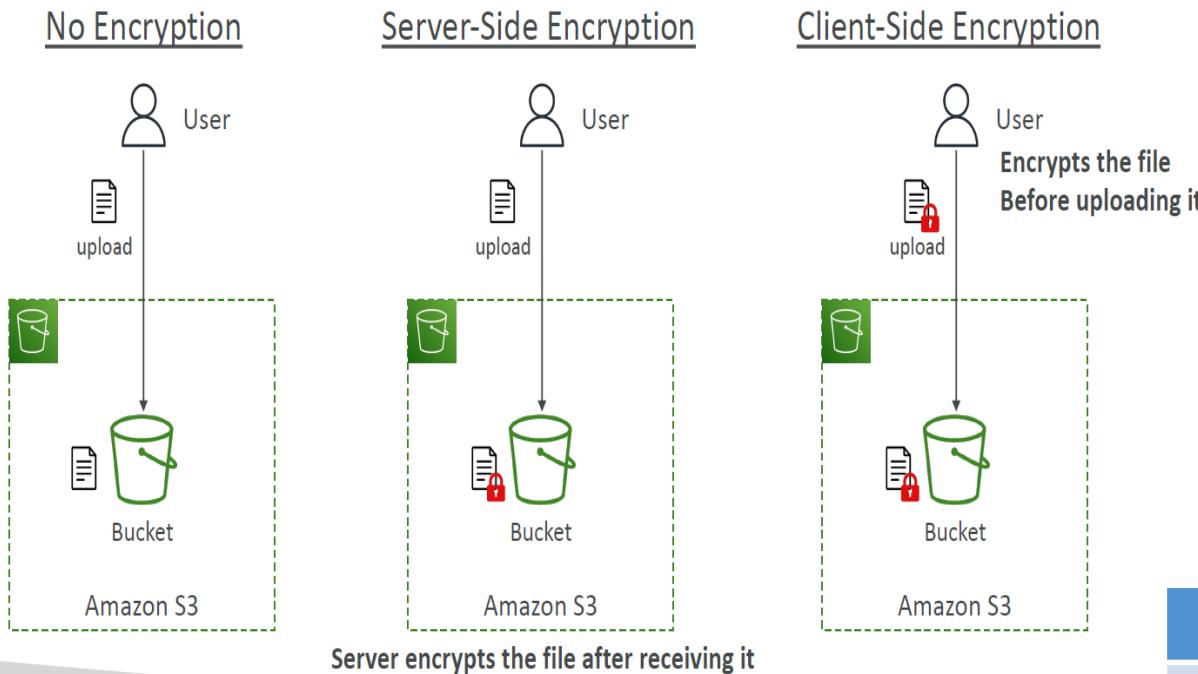
Auto Scaling Strategies include: Manual Scaling, Dynamic Scaling (Simple/Step Scaling, Target Tracking Scaling, Scheduled Scaling), and Predictive Scaling.



Global Accelerator is a good fit for non-HTTP use cases - Global Accelerator is a good fit for non-HTTP use cases, such as gaming (UDP), IoT (MQTT), or Voice over IP, as well as for HTTP use cases that specifically require static IP addresses or deterministic, fast regional failover.

Global Accelerator provides static IP addresses that act as a fixed entry point to your applications - It provides static IP addresses that provide a fixed entry point to your applications and eliminate the complexity of managing specific IP addresses for different AWS Regions and Availability Zones.

S3 Encryption



AWS Snow Family for Data Migrations



	Snowcone	Snowball Edge Storage Optimized	Snowmobile
Storage Capacity	8 TB usable	80 TB usable	< 100 PB
Migration Size	Up to 24 TB, online and offline	Up to petabytes, offline	Up to exabytes, offline
DataSync agent	Pre-installed		
Storage Clustering		Up to 15 nodes	

What is Edge Computing?

- Process data while it's being created on **an edge location**
 - A truck on the road, a ship on the sea, a mining station underground...



- These locations may have
 - Limited / no internet access
 - Limited / no easy access to computing power
- We setup a **Snowball Edge / Snowcone** device to do edge computing
- Use cases of Edge Computing:
 - Preprocess data
 - Machine learning at the edge
 - Transcoding media streams
- Eventually (if need be) we can ship back the device to AWS (for transferring data for example)

AWS Machine Learning - Summary

- **Rekognition:** face detection, labeling, celebrity recognition
- **Transcribe:** audio to text (ex: subtitles)
- **Polly:** text to audio
- **Translate:** translations
- **Lex:** build conversational bots – chatbots
- **Connect:** cloud contact center
- **Comprehend:** natural language processing
- **SageMaker:** machine learning for every developer and data scientist
- **Forecast:** build highly accurate forecasts
- **Kendra:** ML-powered search engine
- **Personalize:** real-time personalized recommendations
- **Textract:** detect text and data in documents

Advanced Identity - Summary

- IAM
 - Identity and Access Management inside your AWS account
 - For users that you trust and belong to your company
- Organizations: manage multiple AWS accounts
- Security Token Service (STS): temporary, limited-privileges credentials to access AWS resources
- Cognito: create a database of users for your mobile & web applications
- Directory Services: integrate Microsoft Active Directory in AWS
- Single Sign-On (SSO): one login for multiple AWS accounts & applications

Monitoring Summary

- **CloudWatch:**
 - **Metrics:** monitor the performance of AWS services and billing metrics
 - **Alarms:** automate notification, perform EC2 action, notify to SNS based on metric
 - **Logs:** collect log files from EC2 instances, servers, Lambda functions...
 - **Events (or EventBridge):** react to events in AWS, or trigger a rule on a schedule
- **CloudTrail:** audit API calls made within your AWS account
- **CloudTrail Insights:** automated analysis of your CloudTrail Events
- **X-Ray:** trace requests made through your distributed applications
- **Service Health Dashboard:** status of all AWS services across all regions
- **Personal Health Dashboard:** AWS events that impact your infrastructure
- **Amazon CodeGuru:** automated code reviews and application performance recommendations

Deployment - Summary

- CloudFormation: (AWS only)
 - Infrastructure as Code, works with almost all of AWS resources
 - Repeat across Regions & Accounts
- Beanstalk: (AWS only)
 - Platform as a Service (PaaS), limited to certain programming languages or Docker
 - Deploy code consistently with a known architecture: ex, ALB + EC2 + RDS
- CodeDeploy (hybrid): deploy & upgrade any application onto servers
- Systems Manager (hybrid): patch, configure and run commands at scale
- OpsWorks (hybrid): managed Chef and Puppet in AWS

Developer Services - Summary

- **CodeCommit**: Store code in private git repository (version controlled)
- **CodeBuild**: Build & test code in AWS
- **CodeDeploy**: Deploy code onto servers
- **CodePipeline**: Orchestration of pipeline (from code to build to deploy)
- **CodeArtifact**: Store software packages / dependencies on AWS
- **CodeStar**: Unified view for allowing developers to do CICD and code
- **Cloud9**: Cloud IDE (Integrated Development Environment) with collab
- **AWS CDK**: Define your cloud infrastructure using a programming language

You need a unified user interface that gives you visibility, control, and patching capabilities for your EC2 Instances on AWS, as well as for servers running in your on-premises data centers. Which service should you use?

Storage Gateway

OpsWorks

Elastic Container Service

Systems Manager

• Global DNS: Route 53

- Great to route users to the closest deployment with least latency
- Great for disaster recovery strategies

• Global Content Delivery Network (CDN): CloudFront

- Replicate part of your application to AWS Edge Locations – decrease latency
- Cache common requests – improved user experience and decreased latency

• S3 Transfer Acceleration

- Accelerate global uploads & downloads into Amazon S3

• AWS Global Accelerator:

- Improve global application availability and performance using the AWS global network

Section Summary: Security & Compliance

- Shared Responsibility on AWS
- Shield: Automatic DDoS Protection + 24/7 support for advanced
- WAF: Firewall to filter incoming requests based on rules
- KMS: Encryption keys managed by AWS
- CloudHSM: Hardware encryption, we manage encryption keys
- AWS Certificate Manager: provision, manage, and deploy SSL/TLS Certificates
- Artifact: Get access to compliance reports such as PCI, ISO, etc...
- GuardDuty: Find malicious behavior with VPC, DNS & CloudTrail Logs
- Inspector: For EC2 only, install agent and find vulnerabilities

Elastic Beanstalk as its meant for automatic application deployment when you upload your code. Then Elastic Beanstalk automatically handles the deployment, from capacity provisioning, load balancing, auto-scaling to application health monitoring.

An IT company would like to move its IT resources (including any data and applications) from an AWS Region in the US to another AWS Region in Europe. Which of the following represents the correct solution for this use-case?

- The company should just start creating new resources in the destination AWS Region and then migrate the relevant data and applications into this new AWS Region** (Correct)

Basic - The basic plan only provides access to the following:

Customer Service & Communities - 24x7 access to customer service, documentation, whitepapers, and support forums. AWS Trusted Advisor - Access to the 7 core Trusted Advisor checks and guidance to provision your resources following best practices to increase performance and improve security. AWS Personal Health Dashboard - A personalized view of the health of AWS services, and alerts when your resources are impacted.

Developer - AWS recommends Developer Support if you are testing or doing early development on AWS and want the ability to get email-based technical support during business hours as well as general architectural guidance as you build and test. This plan provides access to just the 7 core Trusted Advisor checks.

The DevOps team at an IT company is moving 500 GB of data from an EC2 instance to an S3 bucket in the same region. Which of the following scenario captures the correct charges for this data transfer?



The company would not be charged for this data transfer

(Correct)

A company wants to improve the resiliency of its flagship application so it wants to move from its traditional database system to a managed AWS database service to support active-active configuration in both the East and West US AWS regions. The active-active configuration with cross-region support is the prime criteria for any database solution that the company considers.

Which AWS database service is the right fit for this requirement?

- Amazon DynamoDB with global tables** **(Correct)**

A multi-national corporation wants to get expert professional advice on migrating to AWS and managing their applications on AWS Cloud. Which of the following entities would you recommend for this engagement?

APN Consulting Partner

(Correct)

APN Partner Types

APN Consulting Partners

APN Consulting Partners are professional services firms that help customers of all types and sizes design, architect, build, migrate, and manage their workloads and applications on AWS, accelerating their journey to the cloud. APN Consulting Partners often implement Technology Partner solutions in addition to the professional services they offer.

APN Consulting Partners include system integrators, strategic consultancies, agencies, managed service providers, and value-added resellers.

[Learn more »](#)

APN Technology Partners

APN Technology Partners provide hardware, connectivity services, or software solutions that are either hosted on, or integrated with, the AWS Cloud. Technology Partner products are often delivered as components to broader AWS customer solutions and can be delivered globally by Consulting Partners through AWS Marketplace, bundled solutions, or directly from APN Technology Partners.

APN Technology Partners include original equipment manufacturers (OEMs), semiconductor manufacturers, network carriers, SaaS Providers, and independent software vendors (ISVs).

[Learn more »](#)

A startup wants to provision an EC2 instance for the lowest possible cost for a long-term duration but needs to make sure that the instance would never be interrupted. As a Cloud Practitioner, which of the following options would you recommend?



Reserved Instance

(Correct)

On-Demand

With On-Demand instances, you pay for compute capacity by the hour or the second depending on which instances you run. No longer-term commitments or upfront payments are needed. You can increase or decrease your compute capacity depending on the demands of your application and only pay the specified per hourly rates for the instance you use.

On-Demand instances are recommended for:

- Users that prefer the low cost and flexibility of Amazon EC2 without any up-front payment or long-term commitment
- Applications with short-term, spiky, or unpredictable workloads that cannot be interrupted
- Applications being developed or tested on Amazon EC2 for the first time

[See On-Demand pricing »](#)

Savings Plans

Savings Plans are a flexible pricing model that offer low prices on EC2 and Fargate usage, in exchange for a commitment to a consistent amount of usage (measured in \$/hour) for a 1 or 3 year term.

Dedicated Hosts

A Dedicated Host is a physical EC2 server dedicated for your use. Dedicated Hosts can help you reduce costs by allowing you to use your existing server-bound software licenses, including Windows Server, SQL Server, and SUSE Linux Enterprise Server (subject to your license terms), and can also help you meet compliance requirements. [Learn more](#).

- Can be purchased On-Demand (hourly).
- Can be purchased as a Reservation for up to 70% off the On-Demand price.

[See Dedicated pricing »](#)

Spot instances

Amazon EC2 Spot instances allow you to request spare Amazon EC2 computing capacity for up to 90% off the On-Demand price. [Learn More](#).

Spot instances are recommended for:

- Applications that have flexible start and end times
- Applications that are only feasible at very low compute prices
- Users with urgent computing needs for large amounts of additional capacity

[See Spot pricing »](#)

A research group wants to use EC2 instances to run a scientific computation application that has a fault tolerant architecture. The application needs high-performance hardware disks that provide fast I/O performance. As a Cloud Practitioner, which of the following storage options would you recommend as the MOST cost-effective solution?

Reserved Instances

Reserved Instances provide you with a significant discount (up to 75%) compared to On-Demand instance pricing. In addition, when Reserved Instances are assigned to a specific Availability Zone, they provide a capacity reservation, giving you additional confidence in your ability to launch instances when you need them.

For applications that have steady state or predictable usage, Reserved Instances can provide significant savings compared to using On-Demand instances. See [How to Purchase Reserved Instances](#) for more information.

Reserved Instances are recommended for:

- Applications with steady state usage
- Applications that may require reserved capacity
- Customers that can commit to using EC2 over a 1 or 3 year term to reduce their total computing costs

A data analytics company is running a proprietary batch analytics application on AWS and wants to use a storage service which would be accessed by hundreds of EC2 instances simultaneously to append data to existing files. As a Cloud Practitioner, which AWS service would you suggest for this use-case?

"EFS" - Amazon EFS is a file storage service for use with Amazon EC2. **Amazon EFS provides a file system interface, file system access semantics, and concurrently-accessible storage for up to thousands of Amazon EC2 instances.** Amazon EFS uses the Network File System protocol.

What are the different gateway types supported by AWS Storage Gateway service?

Tape Gateway, Object Gateway and Volume Gateway

Tape Gateway, File Gateway and Volume Gateway (Correct)

A social media company wants to protect its web application from common web exploits such as SQL injection and cross-site scripting. Which of the following AWS services can be used to address this use-case?

AWS Web Application Firewall (WAF) (Correct)

Which AWS Route 53 routing policy would you use to route traffic to multiple resources and also choose how much traffic is routed to each resource?

When you create a record, you choose a routing policy, which determines how Amazon Route 53 responds to queries: **Amazon Route 53 is a highly available and scalable cloud Domain Name System (DNS) web service**

- **Simple routing policy** – Use for a single resource that performs a given function for your domain, for example, a web server that serves content for the example.com website.
- **Failover routing policy** – Use when you want to configure active-passive failover.
- **Geolocation routing policy** – Use when you want to route traffic based on the location of your users.
- **Geoproximity routing policy** – Use when you want to route traffic based on the location of your resources and, optionally, shift traffic from resources in one location to resources in another.
- **Latency routing policy** – Use when you have resources in multiple AWS Regions and you want to route traffic to the region that provides the best latency.
- **Multivalue answer routing policy** – Use when you want Route 53 to respond to DNS queries with up to eight healthy records selected at random.
- **Weighted routing policy** – Use to route traffic to multiple resources in proportions that you specify.

A company wants to move to AWS cloud and release new features with quick iterations by utilizing relevant AWS services whenever required. Which of the following characteristics of AWS Cloud does it want to leverage?

Agility

In the world of cloud computing, "Agility" refers to the ability to rapidly develop, test and launch software applications that drive business growth. Another way to explain "Agility" - AWS provides a massive global cloud infrastructure that allows you to quickly innovate, experiment and iterate. Instead of waiting weeks or months for hardware, you can instantly deploy new applications. This ability is called Agility.

Incorrect options:

Elasticity - This refers to the ability to acquire resources as you need and release when they are no longer needed is termed as Elasticity of the Cloud.

Reliability - This refers to the ability of a system to recover from infrastructure or service disruptions, by dynamically acquiring computing resources to meet demand, and mitigate disruptions.

Scalability - Scalability is the measurement of a system's ability to grow to accommodate an increase in demand, or shrink down to a diminishing demand.

Which of the following AWS services support reservations to optimize costs? (Select three)

Amazon EC2 Reserved Instances: You can use Amazon EC2 Reserved Instances to reserve capacity and receive a discount on your instance usage compared to running On-Demand instances.

Amazon DynamoDB Reserved Capacity: If you can predict your need for Amazon DynamoDB read-and-write throughput, Reserved Capacity offers significant savings over the normal price of DynamoDB provisioned throughput capacity.

Amazon ElastiCache Reserved Nodes: Amazon ElastiCache Reserved Nodes give you the option to make a low, one-time payment for each cache node you want to reserve and, in turn, receive a significant discount on the hourly charge for that node.

Amazon RDS RIs: Like Amazon EC2 RIs, Amazon RDS RIs can be purchased using No Upfront, Partial Upfront, or All Upfront terms. All Reserved Instance types are available for Aurora, MySQL, MariaDB, PostgreSQL, Oracle, and SQL Server database engines.

Amazon Redshift Reserved Nodes: If you intend to keep an Amazon Redshift cluster running continuously for a prolonged period, you should consider purchasing reserved-node offerings. These offerings provide significant savings over on-demand pricing, but they require you to reserve compute nodes and commit to paying for those nodes for either a 1- or 3-year duration.

An intern at an IT company provisioned a Linux based On-demand EC2 instance with per-second billing but terminated it within 30 seconds as he wanted to provision another instance type. What is the duration for which the instance would be charged?

60 seconds - There is a one-minute minimum charge for Linux based EC2 instances, so this is the correct option.

Which AWS Support plan provides architectural guidance contextual to your specific use-cases? Which of the following AWS Support plans provides access to Infrastructure Event Management for an additional fee?

Business

<https://aws.amazon.com/premiumsupport/plans/>

(Correct)

Which of the following AWS services has encryption enabled by default?

CloudTrail Logs

(Correct)

Which of the following AWS Support plans provides access to online training with self-paced labs?

A startup wants to migrate its data and applications from the on-premises data center to AWS Cloud. Which of the following options can be used by the startup to help with this migration? (Select two)



Leverage AWS Professional Services to accelerate the infrastructure migration

(Correct)



Utilize AWS Partner Network (APN) to build a custom solution for this infrastructure migration

(Correct)

Which AWS services can be used to decouple components of a microservices based application on AWS Cloud? (Select two)

A company runs an application on a fleet of EC2 instances. The company wants to automate the traditional maintenance job of running timely assessments and checking for OS vulnerabilities. As a Cloud Practitioner, which service will you suggest for this use case?

Amazon Inspector

Amazon Inspector is an automated security assessment service that helps improve the security and compliance of applications deployed on your Amazon EC2 instances. Amazon Inspector automatically assesses applications for **exposure, vulnerabilities, and deviations from best practices**. After performing an assessment, Amazon Inspector produces a detailed list of security findings prioritized by level of severity. These findings can be reviewed directly or as part of detailed assessment reports which are available via the Amazon Inspector console or API.

Incorrect options:

Amazon GuardDuty - Amazon GuardDuty is a **threat detection service that monitors malicious activity and unauthorized behavior to protect your AWS account**. GuardDuty analyzes billions of events across your AWS accounts from AWS CloudTrail (AWS user and API activity in your accounts), Amazon VPC Flow Logs (network traffic data), and DNS Logs (name query patterns). This service is for AWS account level access, not for instance-level management like an EC2. GuardDuty cannot be used to check OS vulnerabilities.

Amazon Macie - Amazon Macie is a fully managed data security and data privacy service that uses machine learning and pattern matching to discover and protect your sensitive data in AWS. **Macie helps identify and alert you to sensitive data, such as personally identifiable information (PII)**. This service is for securing data and has nothing to do with an EC2 security assessment. Macie cannot be used to check OS vulnerabilities.

AWS Shield - AWS **Shield is a managed Distributed Denial of Service (DDoS) protection service that safeguards applications running on AWS**. AWS Shield provides always-on detection and automatic inline mitigations that minimize application downtime and latency, so there is no need to engage AWS Support to benefit from DDoS protection. Shield is general protection against DDoS attacks for all resources in the AWS network, and not an instance-level security assessment service. Shield cannot be used to check OS vulnerabilities.

Which AWS services can be used to facilitate organizational change management, part of the Reliability pillar of AWS Well-Architected Framework? (Select three)

Amazon CloudWatch

(Correct)

AWS CloudTrail

(Correct)

AWS Config

(Correct)



Configuration change occurs in your AWS resources.



AWS Config records and normalizes the changes into a consistent format.



AWS Config automatically evaluates the recorded configurations against the configurations you specify.



Access change history and compliance results using the console or APIs. CloudWatch Events or SNS alert you when changes occur. Deliver change history and snapshot files to your S3 bucket for analysis.



AWS CloudTrail
Track user activity and detect unusual API usage



Capture
Record activity in AWS services as AWS CloudTrail events



Store
AWS CloudTrail delivers events to the AWS CloudTrail console, Amazon S3 buckets, and optionally Amazon CloudWatch Logs



Act
Use Amazon CloudWatch Alarms and Events to take action when important events are detected



Reliability in the cloud - Foundations, Change Management, Failure Management. Being aware of how change affects a system (change management) allows you to plan proactively, and monitoring allows you to quickly identify trends that could lead to capacity issues or SLA breaches.

Amazon CloudWatch - Amazon CloudWatch is a monitoring and observability service built for DevOps engineers, developers, site reliability engineers (SREs), and IT managers. CloudWatch provides data and actionable insights to monitor applications, respond to system-wide performance changes, optimize resource utilization, and get a unified view of operational health.

Which of the following is CORRECT regarding removing an AWS account from AWS Organizations?

You can remove an account from your organization only **if the account has the information that is required for it to operate as a standalone account.** For each account that you want to make standalone, you must accept the AWS Customer Agreement, choose a support plan, provide and verify the required contact information, and provide a current payment method. AWS uses the payment method to charge for any billable (not AWS Free Tier) AWS activity that occurs while the account isn't attached to an organization.

Which security service of AWS is enabled for all AWS customers, by default, at no additional cost?

AWS Shield Standard defends against most common, frequently occurring network and transport layer DDoS attacks that target your website or applications. While AWS Shield Standard helps protect all AWS customers, you get better protection if you are using Amazon CloudFront and Amazon Route 53. All **AWS customers benefit from the automatic protections of AWS Shield Standard, at no additional charge.**

AWS Shield Advanced provides expanded DDoS attack protection for web applications running on which of the following resources? (Select two)

AWS Shield Advanced provides expanded DDoS attack protection for web applications running on the following resources: **Amazon Elastic Compute Cloud, Elastic Load Balancing (ELB), Amazon CloudFront, Amazon Route 53, AWS Global Accelerator.**

Six Advantages of Cloud Computing

[PDF](#) | [RSS](#)

- **Trade capital expense for variable expense** – Instead of having to invest heavily in data centers and servers before you know how you're going to use them, you can pay only when you consume computing resources, and pay only for how much you consume.
- **Benefit from massive economies of scale** – By using cloud computing, you can achieve a lower variable cost than you can get on your own. Because usage from hundreds of thousands of customers is aggregated in the cloud, providers such as AWS can achieve higher economies of scale, which translates into lower pay as-you-go prices.
- **Stop guessing capacity** – Eliminate guessing on your infrastructure capacity needs. When you make a capacity decision prior to deploying an application, you often end up either sitting on expensive idle resources or dealing with limited capacity. With cloud computing, these problems go away. You can access as much or as little capacity as you need, and scale up and down as required with only a few minutes' notice.
- **Increase speed and agility** – In a cloud computing environment, new IT resources are only a click away, which means that you reduce the time to make those resources available to your developers from weeks to just minutes. This results in a dramatic increase in agility for the organization, since the cost and time it takes to experiment and develop is significantly lower.
- **Stop spending money running and maintaining data centers** – Focus on projects that differentiate your business, not the infrastructure. Cloud computing lets you focus on your own customers, rather than on the heavy lifting of racking, stacking, and powering servers.
- **Go global in minutes** – Easily deploy your application in multiple regions around the world with just a few clicks. This means you can provide lower latency and a better experience for your customers at minimal cost.

Which of the following is the MOST cost-effective option to purchase an EC2 Reserved Instance? All you need to remember is that a 3 years term would always be more cost-effective than a 1-year term. Then within a term, "all upfront" is better than "partial upfront" which in turn is better than "no upfront" from a cost savings perspective.

A big data analytics company is moving its IT infrastructure from an on-premises data center to AWS Cloud. The company has some server-bound software licenses that it wants to use on AWS. As a Cloud Practitioner, which of the following EC2 instance types would you recommend to the company?

Differences between Dedicated Hosts and Dedicated Instances

Dedicated Hosts and Dedicated Instances can both be used to launch Amazon EC2 instances onto physical servers that are dedicated for your use.

There are no performance, security, or physical differences between Dedicated Instances and instances on Dedicated Hosts. However, there are some differences between the two. The following table highlights some of the key differences between Dedicated Hosts and Dedicated Instances:

	Dedicated Host	Dedicated Instance
Billing	Per-host billing	Per-instance billing
Visibility of sockets, cores, and host ID	Provides visibility of the number of sockets and physical cores	No visibility
Host and instance affinity	Allows you to consistently deploy your instances to the same physical server over time	Not supported
Targeted instance placement	Provides additional visibility and control over how instances are placed on a physical server	Not supported
Automatic instance recovery	Supported. For more information, see Host recovery .	Supported
Bring Your Own License (BYOL)	Supported	Not supported



As an AWS customer, you want the most value from your investment. Trusted Advisor can help.



AWS Trusted Advisor

Trusted Advisor scans your AWS infrastructure, compares it to AWS best practices in five categories, and provides recommended actions



Cost Optimization



investigation recommended



Performance



no problem detected



Security



action recommended



Fault Tolerance



no problem detected



Service limits



no problem detected

Recommendations & action links

An IT company wants to run a log backup process every Monday at 2 AM. The usual runtime of the process is 5 minutes. As a Cloud Practitioner, which AWS services would you recommend to build a serverless solution for this use-case? (Select two)

Lambda (Correct)

CloudWatch (Correct)

A company's flagship application runs on a fleet of Amazon EC2 instances. As per the new policies, the system administrators are looking for the best way to provide secure shell access to AWS EC2 instances without opening new ports or using public IP addresses.

Which tool/service will help you achieve this requirement?

AWS Systems Manager Session Manager (Correct)

An online gaming company wants to block users from certain geographies from accessing its content. Which AWS services can be used to accomplish this task? (Select two) **AWS WAF, Route 53**

AWS Budgets - Cost budget, Usage budget, Reservation budget and Savings Plans budget.

AWS Cost Explorer Features

Get started quickly

A set of default reports are included to help you quickly gain insight into your cost drivers and usage trends.

Forecast future costs and usage

Use forecasting to get a better idea of what your costs and usage may look like in the future, so that you can plan ahead.

Set time interval and granularity

Set a custom time period, and determine whether you would like to view your data at a monthly or daily level of granularity.

Save your progress

Once you arrive at a helpful view, save your progress as a new report that you can refer back to in the future.

Filter/Group your data

Dig deeper into your data by taking advantage of filtering and grouping functionality, using a variety of available dimensions.

Build custom applications

Directly access the interactive, ad-hoc analytics engine that powers AWS Cost Explorer.

Bob and Susan each have an AWS account in AWS Organizations. Susan has five Reserved Instances (RIs) of the same type and Bob has none. During one particular hour, Susan uses three instances and Bob uses six for a total of nine instances on the organization's consolidated bill.

Which of the following statements are correct about consolidated billing in AWS Organizations? (Select two)

AWS bills five instances as Reserved Instances, and the remaining four instances as regular instances (Correct)

Bob receives the cost-benefit from Susan's Reserved Instances only if he launches his instances in the same Availability Zone where Susan purchased her Reserved Instances (Correct)

AWS Well-Architected Framework

Foundations are part of the **Reliability** pillar of the AWS Well-Architected Framework.

AWS states that before architecting any system, foundational requirements that influence reliability should be in place. The services that are part of foundations are: **Amazon VPC, AWS Trusted Advisor, AWS Service Quotas (formerly called AWS Service Limits).**

- 1) Operational Excellence
- 2) Security
- 3) Reliability
- 4) Performance Efficiency
- 5) Cost Optimization
- 6) Sustainability

1) Who is responsible for configuration management under the AWS shared responsibility model?

- A. It is solely the responsibility of the customer.
- B. It is solely the responsibility of AWS.
- C. It is shared between AWS and the customer.
- D. It is not part of the AWS shared responsibility model.



Correct Answer: C

AWS maintains the configuration of its infrastructure devices, but a customer is responsible for configuring their own guest operating systems, databases, and applications.

2) Which compute hosting model should be accounted for in the Total Cost of Ownership (TCO) when undertaking a cost analysis that allows physical isolation of a customer workload?

- 
- A. Dedicated Hosts
 - B. Reserved Instances
 - C. On-Demand Instances
 - D. No Upfront Reserved Instances

Correct Answer: A

Use Dedicated Hosts to launch Amazon EC2 instances on physical servers that are dedicated for your use. Dedicated Hosts give you additional visibility and control over how instances are placed on a physical server, and you can reliably use the same physical server over time. As a result, Dedicated Hosts enable you to use your existing server-bound software licenses like Windows Server and address corporate compliance and regulatory requirements.

1) Who is responsible for configuration management under the AWS shared responsibility model?

- A. It is solely the responsibility of the customer.
- B. It is solely the responsibility of AWS.
- C. It is shared between AWS and the customer.
- D. It is not part of the AWS shared responsibility model.



Correct Answer: C

AWS maintains the configuration of its infrastructure devices, but a customer is responsible for configuring their own guest operating systems, databases, and applications.

12) Which pillar of the AWS Well-Architected Framework is supported by the design philosophy of performing operations as code?

- A. Performance efficiency
 - B. Operational excellence
 - C. Reliability
 - D. Security
- 

Correct Answer: B

1. Operational Excellence

The Operational Excellence pillar includes the ability to support development and run workloads effectively, gain insight into their operation, and continuously improve, supporting processes and procedures to delivery business value. You can find prescriptive guidance on implementation in the [Operational Excellence Pillar whitepaper](#).

Design Principles

There are five design principles for operational excellence in the cloud:

- Perform operations as code
- Make frequent, small, reversible changes
- Refine operations procedures frequently
- Anticipate failure
- Learn from all operational failures

14) Which of the following operational controls do users completely inherit from AWS as part of the AWS shared responsibility model?

- •A. Security management of data center
•B. Patch management
•C. Configuration management
•D. User and access management

Correct Answer: A

Security and Compliance is a shared responsibility between AWS and the customer. This shared model can help relieve the customer's operational burden as AWS operates, manages and controls the components from the host operating system and virtualization layer down to the physical security of the facilities in which the service operates. The customer assumes responsibility and management of the guest operating system (including updates and security patches), other associated application software as well as the configuration of the AWS provided security group firewall. Customers should carefully consider the services they choose as their responsibilities vary depending on the services used, the integration of those services into their IT environment, and applicable laws and regulations. The nature of this shared responsibility also provides the flexibility and customer control that permits the deployment. As shown in the chart below, this differentiation of responsibility is commonly referred to as Security "of" the Cloud versus Security "in" the Cloud.

16) Which pillar of the AWS Well-Architected Framework is designed on the idea of frequent, minor, reversible changes?

- 
- A. Reliability
 - B. Operational excellence
 - C. Performance efficiency
 - D. Cost optimization

Correct Answer: B

There are five design principles for operational excellence in the cloud:

- 1) **Perform operations as code**
- 2) **Make frequent, small, reversible changes**
- 3) **Refine operations procedures frequently**
- 4) **Anticipate failure**
- 5) **Learn from all operational failures**

17) Which AWS service can identify the person who made the API request when an Amazon EC2 instance is terminated?

- 
- A. Amazon CloudWatch
 - B. AWS CloudTrail
 - C. AWS X-Ray
 - D. AWS Identity and Access Management (IAM)

Correct Answer: B

AWS CloudTrail allows AWS customers to **record API calls, sending log files to Amazon S3 buckets for storage.**

20) Permissions for which of the following are managed by service control policies (SCPs)?

- A. Availability Zones
- B. AWS Regions
- •C. AWS Organizations
- D. Edge locations

Correct Answer: C

Service Control Policies (SCPs) are IAM-like policies to manage permissions in AWS Organizations. SCPs restrict the actions allowed for accounts within the organization making each one of them compliant with your guidelines.

21) Which of the following AWS services should a client utilize to audit the change management of AWS resources?

- 
- A. AWS Config
 - B. AWS Trusted Advisor
 - C. Amazon CloudWatch
 - D. Amazon Inspector

Correct Answer: A

AWS Config is a service that enables you to assess, audit, and evaluate the configurations of your AWS resources. Config continuously monitors and records your AWS resource configurations and allows you to automate the evaluation of recorded configurations against desired configurations. With Config, you can review changes in configurations and relationships between AWS resources, dive into detailed resource configuration histories, and determine your overall compliance against the configurations specified in your internal guidelines. This enables you to simplify compliance auditing, security analysis, change management, and operational troubleshooting.

22) What is raised when a business deploys web servers across several AWS Regions?

- A. Coupling
- B. Availability
- C. Security
- D. Durability



Correct Answer: D

AWS measures durability as a percentage. For example, the S3 Standard Tier is designed for **99.99999999% durability**.

23) Which of the following is a shared control between a client and AWS under the shared responsibility model?

- A. Physical controls
- B. Patch management
- C. Zone security
- D. Data center auditing

Correct Answer: B

Patch Management: AWS is responsible for patching and fixing flaws within the infrastructure, but customers are responsible for patching their guest OS and applications

24) According to the AWS shared responsibility model, which job is shared between AWS and the customer?

- A. Physical and environmental controls
- B. Server hardware management and encryption
- C. Application security
- D. Patch management and configuration management



Correct Answer: D

Shared Controls which apply to both the infrastructure layer and customer layers, but in completely separate contexts or perspectives. In a shared control, AWS provides the requirements for the infrastructure and the customer must provide their own control implementation within their use of AWS services.

Examples include:

- ☞ Patch Management AWS is responsible for patching and fixing flaws within the infrastructure, but customers are responsible for patching their guest OS and applications.
- ☞ Configuration Management AWS maintains the configuration of its infrastructure devices, but a customer is responsible for configuring their own guest operating systems, databases, and applications.

25) How should an application be created to function on the AWS Cloud in accordance with best practices?

- A. Use tightly coupled components.
- •B. Use loosely coupled components.
- C. Use infrequently coupled components.
- D. Use frequently coupled components.

Correct Answer: B

The more **loosely coupled** the application **components** are, the better they will scale.

26) Which cloud architectural concept is supported by a system that can scale in terms of users, traffic, or data quantity without sacrificing performance?

- A. Think parallel
- •B. Implement elasticity
- C. Decouple your components
- D. Design for failure

Correct Answer: B

Elasticity - AWS Well-Architected Framework Elasticity

The ability to acquire resources as you need them and release resources when you no longer need them. In the cloud, you want to do this automatically.

27) The following describes an application that spans various Availability Zones:

- •A. being highly available
•B. having global reach
•C. using an economy of scale
•D. having elasticity

Correct Answer: A

An **Availability Zone** (AZ) is one or more discrete data centers with redundant power, networking, and connectivity in an **AWS** Region. AZs give customers the ability to operate production applications and databases that are more highly available, fault tolerant, and scalable than would be possible from a single data center.

28) Which duty is the customer's responsibility while administering AWS Lambda functions under the AWS shared responsibility model?

- 
- A. Creating versions of Lambda functions
 - B. Maintaining server and operating systems
 - C. Scaling Lambda resources according to demand
 - D. Updating the Lambda runtime environment

Correct Answer: A

You can use versions to manage the deployment of your functions. For example, you can publish a new version of a function for beta testing without affecting users of the stable production version. Lambda creates a new version of your function each time that you publish the function. The new version is a copy of the unpublished version of the function.

29) Which of the following is not a duty of the client under the AWS shared responsibility model? (Select two.)

- A. Decommissioning of physical storage devices
- B. Security group and ACL configuration
- C. Patch management of an Amazon RDS instance operating system
- D. Controlling physical access to data centers
- E. Patch management of an Amazon EC2 instance operating system

Correct Answer: AD

30) Which service or functionality does AWS Enterprise Support provide that other AWS Support subscriptions do not?

- A. AWS Trusted Advisor
- B. AWS Support case
- C. Concierge team
- D. Amazon Connect



Correct Answer: C

- AWS Support Concierge Included as part of the Enterprise Support plan, the Support Concierge Team are **AWS billing and account experts** that specialize in working with enterprise accounts.

31) What attributes of an AWS account can AWS Trusted Advisor monitor and advise on? (Select two.)

- •A. Compliance with security best practices
- B. Application performance
- C. Network utilization
- •D. Cost optimization
- E. Compliance status

Correct Answer: AD

Cost Optimization



0 ✓ 9 ▲ 0 !

\$7,516.85

Potential monthly savings

Performance



3 ✓ 7 ▲ 0 !

Security



2 ✓ 4 ▲ 11 !

Fault Tolerance



0 ✓ 15 ▲ 5 !

Service Limits



37 ✓ 0 ▲ 1 !

32) Which feature enables Amazon EC2 instances to be more elastic in response to changing workload demand?

- 
- A. Virtualization Management
 - B. Hardware management
 - C. Encryption management
 - D. Facilities management
 - E. Firewall management

Correct Answer: A

33) What is the unique advantage that Enterprise Support customers receive?

- A. Access to a Technical Project Manager
- •B. Access to a Technical Account Manager
- C. Access to a Cloud Support Engineer
- D. Access to a Solutions Architect

Correct Answer: B

Technical account management

Designated point of contact

A Technical Account Manager (TAM) is your designated technical point of contact who helps you onboard, provides advocacy and guidance to help plan and build solutions using best practices, coordinates access to subject matter experts, assists with case management, presents insights and recommendations on your AWS spend, workload optimization, and event management, and proactively keeps your AWS environment healthy.

Access to subject-matter experts

Cloud Support Engineers, Solutions Architects, and product teams are available to provide guidance and help as needed. The AWS Trust & Safety team assists you when your AWS resources are used to engage in abusive behaviors, such as spam, port scanning, denial-of-service (DoS) attacks, or malware.

34) What are the advantages of AWS Cloud elasticity?

- A. Ensure web traffic is automatically spread across multiple AWS Regions.
- B. Minimize storage costs by automatically archiving log data.
- C. Enable AWS to automatically select the most cost-effective services.
- D. Automatically adjust the required compute capacity to maintain consistent performance.



Correct Answer: D

Elasticity

The ability to acquire resources as you need them and release resources when you no longer need them. In the cloud, you want to do this automatically.

35) Which cost components should be addressed when developing a cloud Total Cost of Ownership (TCO) model for AWS workloads? (Select three.)

- •A. Compute costs
- •B. Facilities costs
- •C. Storage costs
- •D. Data transfer costs
- •E. Network infrastructure costs
- •F. Hardware lifecycle costs

Correct Answer: ACE

Components of Total Cost for Technology

- Hardware acquisition
- Software acquisition
- Infrastructure
- Downtime
- Installation
- Maintenance
- Training
- Support
- Space
- Electricity

36) Which of the following are capabilities of the AWS Cost Management tools? (Select two.)

- A. Terminate all AWS resources automatically if budget thresholds are exceeded.
- •B. Break down AWS costs by day, service, and linked AWS account.
- •C. Create budgets and receive notifications if current or forecasted usage exceeds the budgets.
- D. Switch automatically to Reserved Instances or Spot Instances, whichever is most cost-effective.
- E. Move data stored in Amazon S3 to a more cost-effective storage class.

Correct Answer: BC

AWS offers the following tools for Cost Management:

- Billing Dashboard
- Cost Explorer
- Cost Usage and Report

37) How should an Amazon EC2 instance be granted access to an Amazon S3 bucket in accordance with security best practices?

- A. Hard code an IAM user's secret key and access key directly in the application, and upload the file.
- B. Store the IAM user's secret key and access key in a text file on the EC2 instance, read the keys, then upload the file.
- C. Have the EC2 instance assume a role to obtain the privileges to upload the file.
- D. Modify the S3 bucket policy so that any service can upload to it at any time.

Correct Answer: C

38) Which of the following price options would result in the highest reduction after making an Amazon EC2 Dedicated Host reservation?

- A. No upfront payment
- B. Hourly on-demand payment
- C. Partial upfront payment
- ▶ •D. All upfront payment

Correct Answer: D

Reference:

<https://aws.amazon.com/ec2/pricing/reserved-instances/pricing/>

39) Which AWS services and/or features increase availability and mitigate the effect of failures while constructing a typical three-tier web application? (Select two.)

- A. AWS Auto Scaling for Amazon EC2 instances
- B. Amazon VPC subnet ACLs to check the health of a service
- C. Distributed resources across multiple Availability Zones
- D. AWS Server Migration Service (AWS SMS) to move Amazon EC2 instances into a different Region
- E. Distributed resources across multiple AWS points of presence

Correct Answer: AC

You can use Auto Scaling to help make sure that you are running the desired number of healthy EC2 instances across multiple Availability Zones. Auto Scaling can also automatically increase the number of EC2 instances during demand spikes to maintain performance and decrease capacity during less busy periods to optimize costs

40) Which of the following IT responsibilities may AWS relieve a company's IT team of? (Select two.)

- ➡ A. Patching database software
- ➡ B. Storage capacity planning
- C. Creating database schemas
- D. Setting up access controls for data
- E. Writing application code

Correct Answer: AB

41) Which of the following is included in the AWS Trusted Advisor checks? (Select two.)

- ➡ A. Information on Amazon S3 bucket permissions
- ➡ B. AWS service outages
- ➡ C. Multi-factor authentication enabled on the AWS account root user
- D. Available software patches
- E. Number of users in the account

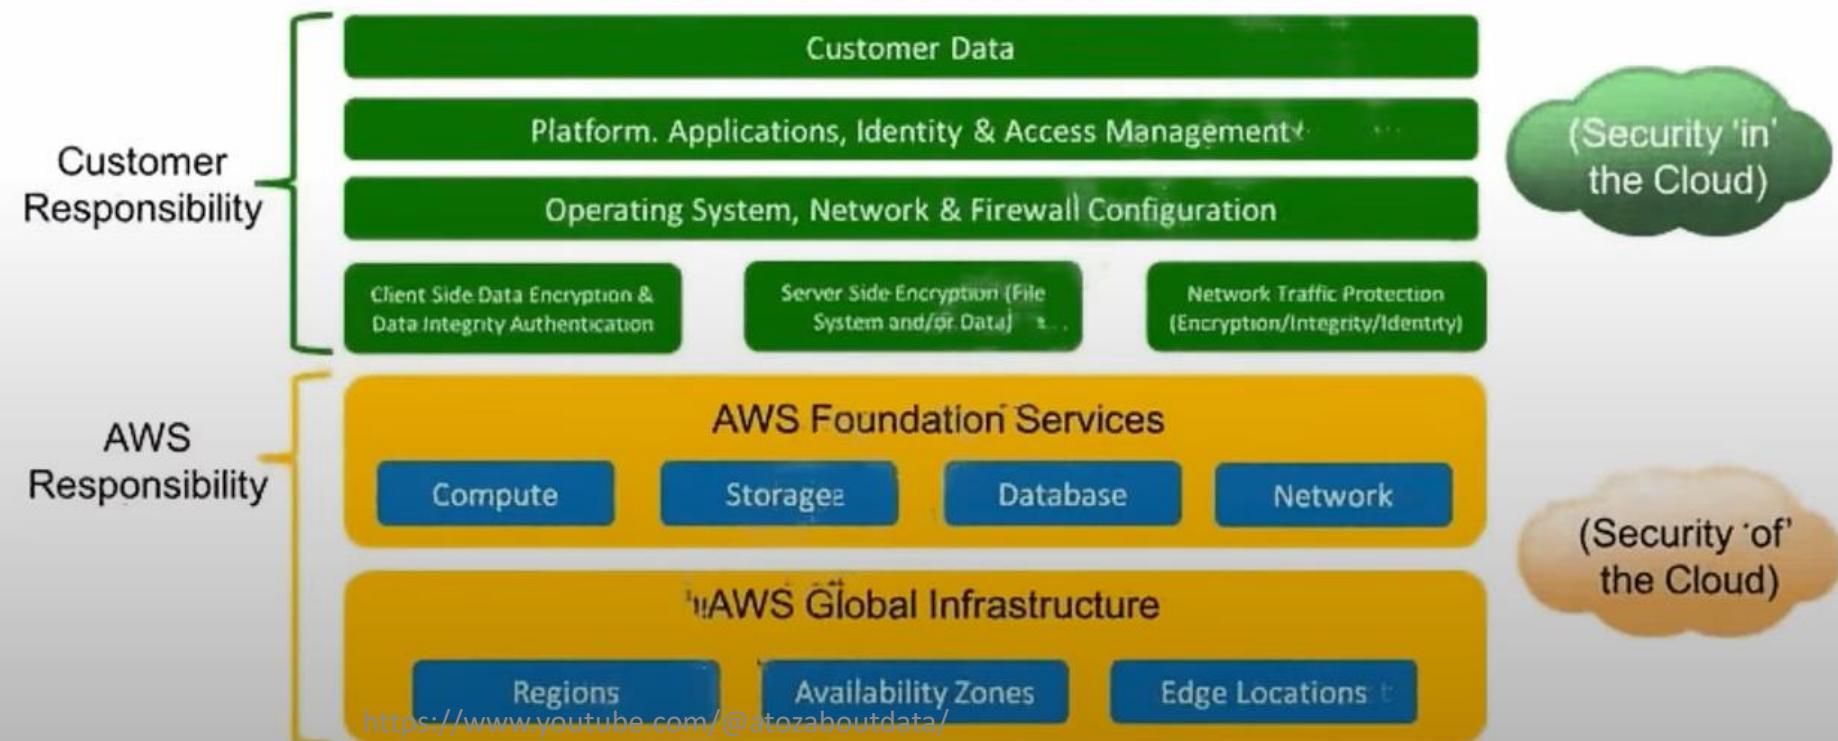
Correct Answer: AC

42) What is AWS accountable for under the AWS shared responsibility model?

- A. Configuring Amazon VPC
- B. Managing application code
- C. Maintaining application traffic
- D. Managing the network infrastructure

Correct Answer: D

Shared Responsibility Model: Infrastructure Services



43) Which of the following tasks may be used to improve Amazon EC2 costs? (Select two.)

- •A. Implementing Auto Scaling groups to add and remove instances based on demand.
- B. Creating a policy to restrict IAM users from creating new instances.
- C. Setting a budget to limit spending on EC2 instances using AWS Budgets.
- •D. Purchasing Reserved Instances.
- E. Adding EC2 instances to a second AWS Region that is geographically close to the end users.

Correct Answer: AD

43) Which of the following tasks may be used to improve Amazon EC2 costs? (Select two.)

- A. Implementing Auto Scaling groups to add and remove instances based on demand.
- B. Creating a policy to restrict IAM users from creating new instances.
- C. Setting a budget to limit spending on EC2 instances using AWS Budgets.
- D. Purchasing Reserved Instances.
- E. Adding EC2 instances to a second AWS Region that is geographically close to the end users.

Correct Answer: AD

44) Which of the following is a customer duty under the AWS shared responsibility model?

- A. Installing security patches for the Xen and KVM hypervisors
- B. Installing operating system patches for Amazon DynamoDB
- C. Installing operating system security patches for Amazon EC2 database instances
- D. Installing operating system security patches for Amazon RDS database instances

Correct Answer: C

45) A firm wishes to develop a new line of business application.

Which design concepts should be applied in accordance with the AWS Well-Architected Framework? (Select two.)

- A. Consolidate multiple AWS accounts into a single account.
- B. Buy and host hardware in the AWS Cloud.
- •C. Decouple the AWS Cloud architecture to break up monolithic deployments.
- D. Move on-premises network hardware to VPCs.
- •E. Design elasticity into the AWS Cloud design.

Correct Answer: CE

46) Which scenarios should a user report to the AWS Abuse team? (Select two.)

- •A. A DDoS attack is being made on an AWS resource.
- B. A SQL injection attack is being made from an IP address that is not an AWS address.
- •C. AWS resources are being used to host objectionable or illegal content.
- D. A company's resources are being used in a way that is inconsistent with corporate policy.
- E. A company is receiving HTTPS requests on a web server that is serving HTTP.

Correct Answer: AC

47) Which of the following is the customer's obligation under the AWS shared responsibility model?

- 
- A. Patching guest OS and applications
 - B. Patching and fixing flaws in the infrastructure
 - C. Physical and environmental controls
 - D. Configuration of AWS infrastructure devices

Correct Answer: A

48) Which AWS serverless platform services are included?

- A. Amazon EC2, Amazon S3, Amazon Athena
- B. Amazon Kinesis, Amazon SQS, Amazon EMR
- C. AWS Step Functions, Amazon DynamoDB, Amazon SNS
- D. Amazon Athena, Amazon Cognito, Amazon EC2



Correct Answer: C

AWS provides a set of fully managed services that you can use to build and run serverless applications. Serverless applications don't require provisioning, maintaining, and administering servers for backend components such as compute, databases, storage, stream processing, message queueing, and more. You also no longer need to worry about ensuring application fault tolerance and availability. Instead, AWS handles all of these capabilities for you. Serverless platform includes: AWS lambda, Amazon S3, DynamoDB, API gateway, Amazon SNS, AWS step functions, Amazon kinesis and developing tools and services.

49) According to the AWS shared responsibility model, which of the following is the customer's obligation?

- A. Patching underlying infrastructure
- B. Physical security
- C. Patching Amazon EC2 instances
- D. Patching network infrastructure

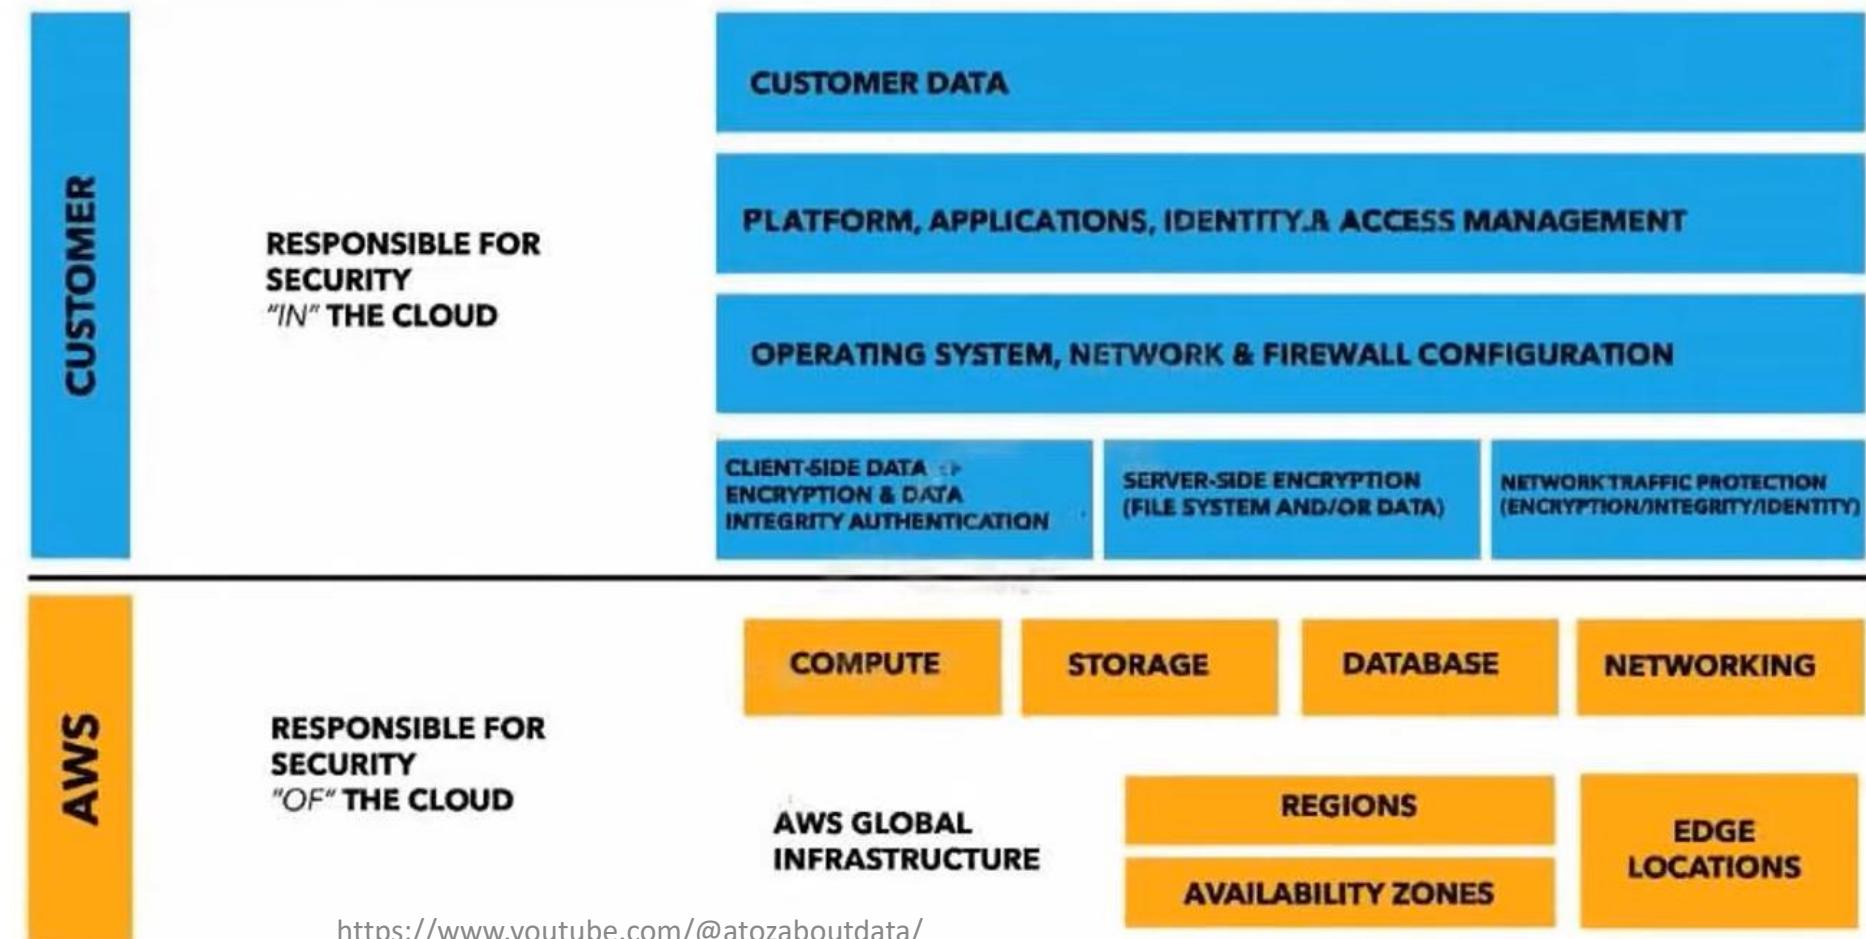


Correct Answer: C

50) Which of the following is the customer's obligation under the shared responsibility model?

- A. Ensuring that disk drives are wiped after use.
- B. Ensuring that firmware is updated on hardware devices.
- C. Ensuring that data is encrypted at rest.
- D. Ensuring that network cables are category six or higher.

Correct Answer: C



51) What is the AWS Cloud's value proposition?

- 
- A. AWS is responsible for security in the AWS Cloud
 - B. No long-term contract is required
 - C. Provision new servers in days
 - D. AWS manages user applications in the AWS Cloud

Correct Answer: B

52) Which of the following describes elasticity in the AWS Cloud? (Select two.)

- A. How quickly an Amazon EC2 instance can be restarted
- •B. The ability to rightsize resources as demand shifts
- C. The maximum amount of RAM an Amazon EC2 instance can use
- D. The pay-as-you-go billing model
- •E. How easily resources can be produced when they are needed

Correct Answer: BE

53) After studying the AWS Cost and Usage report in the AWS Management Console, a Cloud Practitioner detects a billing problem.

Which course of action should be done to rectify this?

- A. Open a detailed case related to billing and submit it to AWS Support for help.
- B. Upload data describing the issue to a new object in a private Amazon S3 bucket.
- C. Create a pricing application and deploy it to a right-sized Amazon EC2 instance for more information.
- D. Proceed with creating a new dashboard in Amazon QuickSight.

Q54 54) Which action, in terms of AWS Cloud architecture, is compatible with the concept of least privilege?

- A. Allow users the minimum access that is needed to do a task.
- B. Prevent managers from accessing important source code.
- C. Provide only the permissions that are needed for users to do their jobs in the current month.
- D. Assign permissions that are based on job titles.

54) Which action, in terms of AWS Cloud architecture, is compatible with the concept of least privilege?

- •A. Allow users the minimum access that is needed to do a task.
•B. Prevent managers from accessing important source code.
•C. Provide only the permissions that are needed for users to do their jobs in the current month.
•D. Assign permissions that are based on job titles.

Correct Answer: A

55) Which of the following are customer duties under the AWS shared responsibility model? (Select two.)

- •A. Setting up server-side encryption on an Amazon S3 bucket
- •B. Amazon RDS instance patching
- C. Network and firewall configurations
- D. Physical security of data center facilities
- E. Compute capacity availability

Correct Answer: AC

56) Which of the following tasks is the AWS customer's duty under the shared responsibility model? (Select two.)

- •A. Ensuring that application data is encrypted at rest
- •B. Ensuring that AWS NTP servers are set to the correct time
- •C. Ensuring that users have received security training in the use of AWS services
- D. Ensuring that access to data centers is restricted
- E. Ensuring that hardware is disposed of properly

Correct Answer: AC

57) Before transferring an environment to the AWS Cloud, a cloud practitioner must receive AWS compliance reports.

How are these reports produced?

- A. Contact the AWS Compliance team
- B. Download the reports from AWS Artifact
- C. Open a case with AWS Support
- D. Generate the reports with Amazon Macie

Correct Answer: B

58) Which change management actions, according to the AWS Well-Architected Framework, should be followed to ensure AWS Cloud reliability? (Select two.)

- •A. Use AWS Config to generate an inventory of AWS resources
- •B. Use service limits to prevent users from creating or making changes to AWS resources
- •C. Use AWS CloudTrail to record AWS API calls into an auditable log file
- •D. Use AWS Certificate Manager to whitelist approved AWS resources and services
- •E. Use Amazon GuardDuty to validate configuration changes made to AWS resources

Correct Answer: AC

59) What may be associated to an Amazon EC2 instance through AWS Identity and Access Management (IAM) to initiate service requests?

- A. Group
- B. Role
- C. Policy
- D. Access key

Correct Answer: B

IAM role management

Q: What is an IAM role?

An IAM role is an IAM entity that defines a set of [permissions](#) for making AWS service requests. IAM roles are not associated with a specific user or group. Instead, trusted entities assume roles, such as IAM users, applications, or AWS services such as EC2.

Q: What problems do IAM roles solve?

IAM roles allow you to delegate access with defined permissions to trusted entities without having to share long-term access keys. You can use IAM roles to delegate access to IAM users managed within your account, to IAM users under a different AWS account, or to an AWS service such as EC2.

Q: How do I get started with IAM roles?

You create a role in a way similar to how you create a user—name the role and attach a policy to it.

For details, see [Creating IAM Roles](#)
<https://www.youtube.com/@atozaboutdata/>

60) Which of the following actions is the customer's responsibility under the AWS shared responsibility model? (Select two.)

- A. Patching operating system components for Amazon Relational Database Server (Amazon RDS)
- B. Encrypting data on the client-side
- C. Training the data center staff
- D. Configuring Network Access Control Lists (ACL)
- E. Maintaining environmental controls within a data center

Correct Answer: BD

61) Which AWS services or functionalities enable failover across several AWS Regions to ensure high availability and low latency? (Select two.)

- •A. Amazon Route 53
- •B. Network Load Balancer
- •C. Amazon S3 Transfer Acceleration
- •D. AWS Global Accelerator
- •E. Application Load Balancer

Correct Answer: AD

62) Which of the following is a duty of AWS under the AWS shared responsibility model?

- A. Enabling server-side encryption for objects stored in S3
- B. Applying AWS IAM security policies
- C. Patching the operating system on an Amazon EC2 instance
- ▶ •D. Applying updates to the hypervisor

Correct Answer: D

63) A copy of the AWS SOC 2 report is required during a compliance review by one of the auditors.

Which service should this request be sent through?

- A. AWS Personal Health Dashboard
- B. AWS Trusted Advisor
- C. AWS Artifact
- D. Amazon S3

Correct Answer: C

AWS Artifact provides on-demand downloads of AWS security and compliance documents, such as AWS ISO certifications, Payment Card Industry (PCI), and Service Organization Control (SOC) reports. You can submit the security and compliance documents (also known as *audit artifacts*) to your auditors or regulators to demonstrate the security and compliance of the AWS infrastructure and services that you use.

64) For 3 hours, 5 minutes, and 6 seconds, a client operates an Amazon Linux EC2 instance on-demand.

How long will the consumer be charged?

- A. 3 hours, 5 minutes
- ▶ •B. 3 hours, 5 minutes, and 6 seconds
- C. 3 hours, 6 minutes
- D. 4 hours

Correct Answer: B

Today, customers may have use cases that require large number of instances running for irregular periods of time, such as dev/test, data processing, analytics, batch processing, image rendering and gaming applications. To improve flexibility and allow customers to focus on their application instead of maximizing usage to the hour, with EC2 per second billing, customers can start, stop, and terminate EC2 instances and take comfort in the fact that their EC2 resources will not be charged for the minutes and seconds remaining in the hours their instances were running.

This change is effective in all AWS Regions and is effective immediately, for Linux instances that are newly launched or already running. Per-second billing is not currently applicable to instances running Microsoft Windows or Linux distributions that have a separate hourly charge.

List prices and Spot Market prices are still listed on a per-hour basis, but bills are calculated down to the second. With Reserved Instances usage you can launch, use, and terminate multiple instances within an hour and get the Reserved Instance Benefit for all of the instances. Bills will now show times in decimal form instead of in minutes and seconds.

65) According to the AWS shared responsibility model, which of the following is exclusively the responsibility of AWS?

- A. Patching of the guest operating system
- B. Security awareness and training
- C. Physical and environmental controls
- D. Development of an IAM password policy



Correct Answer: C

66) What can users do using AWS Marketplace? (Select two.)

- A. Sell unused Amazon EC2 Spot Instances.
- B. Sell solutions to other AWS users.
- C. Buy third-party software that runs on AWS.
- D. Purchase AWS security and compliance documents.
- E. Order AWS Snowball.

Correct Answer: BC

The AWS Marketplace enables qualified partners to market and sell their software to AWS Customers. AWS Marketplace is an online software store that helps customers find, buy, and immediately start using the software and services that run on AWS.

67) What are the possible uses for AWS edge locations? (Select two.)

- A. Hosting applications
- B. Delivering content closer to users
- C. Running NoSQL database caching services
- D. Reducing traffic on the server by caching responses
- E. Sending notification messages to end users

Correct Answer: BD

CloudFront delivers your content through a worldwide network of data centers called edge locations. When a user requests content that you're serving with CloudFront, the user is routed to the edge location that provides the lowest latency (time delay), so that content is delivered with the best possible performance.

68) Amazon Elastic Container Service (Amazon ECS) and Amazon DynamoDB are used by a firm to execute its mission-critical web application. Multiple times during the day, the workload jumps to up to ten times the regular level.

Which AWS Cloud feature helps the business to adapt to these demand changes?

- A. Agility
- B. Global reach
- C. Scalability
- D. Security



Correct Answer: C

69) Which of the following tasks is the customer's duty under the shared responsibility model? (Select two.)

- A. Maintaining the underlying Amazon EC2 hardware.
- B. Managing the VPC network access control lists.
- C. Encrypting data in transit and at rest.
- D. Replacing failed hard disk drives.
- E. Deploying hardware in different Availability Zones.

Correct Answer: BC

The hardware related jobs is the prime responsibility of AWS. VPC network access control lists is something a customer has to do himself to secure the applications. Encrypting data in transit and at rest is a shared responsibility in which AWS plays a part. All hardware related jobs have nothing to do with the customer.

70) Which of the following are AWS obligations, according to the AWS shared responsibility model? (Select two.)

- •A. Network infrastructure and virtualization of infrastructure
- •B. Security of application data
- •C. Guest operating systems
- •D. Physical security of hardware
- •E. Credentials and policies

Correct Answer: AD

71) What charges are included in the comparison of AWS vs on-premises Total Cost of Ownership (TCO)?

- 
- A. Data center security
 - B. Business analysis
 - C. Project management
 - D. Operating system administration

Correct Answer: A

72) Which AWS service supports MySQL and PostgreSQL as relational databases?

- A. Amazon Redshift
- B. Amazon DynamoDB
- C. Amazon Aurora
- D. Amazon Neptune

Correct Answer: C

Amazon Aurora is a MySQL and PostgreSQL-compatible [relational database](#) built for the cloud that combines the performance and availability of traditional enterprise databases with the simplicity and cost-effectiveness of open source databases.

Amazon Aurora is up to five times faster than standard MySQL databases and three times faster than standard PostgreSQL databases. It provides the security, availability, and reliability of commercial databases at 1/10th the cost. Amazon Aurora is fully managed by Amazon Relational Database Service (RDS), which automates time-consuming administration tasks like hardware provisioning, database setup, patching, and backups.

Amazon Aurora features a distributed, fault-tolerant, self-healing storage system that auto-scales up to 128TB per database instance. It delivers high performance and availability with up to 15 low-latency read replicas, point-in-time recovery, continuous backup to [Amazon S3](#), and replication across three Availability Zones.

Visit the [Amazon RDS Management Console](#) to create your first Aurora database instance and start migrating your MySQL and PostgreSQL databases. <https://www.youtube.com/@atozaboutdata/>

73) Which AWS services is natively supported by AWS Snowball Edge?

- A. AWS Server Migration Service (AWS SMS)
- B. Amazon Aurora
- C. AWS Trusted Advisor
- D. Amazon EC2



a. AWS Snowball Edge:

It is a physical device sent by AWS for sending large data.

Space: 80 GB or 40 GB (Two variants)



Snowball Edge (for data transfers)



- Physical data transport solution: move TBs or PBs of data in or out of AWS
- Alternative to moving data over the network (and paying network fees)
- Pay per data transfer job
- Provide block storage and Amazon S3-compatible object storage
- Snowball Edge Storage Optimized
 - 80 TB of HDD capacity for block volume and S3 compatible object storage
- Snowball Edge Compute Optimized
 - 42 TB of HDD capacity for block volume and S3 compatible object storage
- Use cases: large data cloud migrations, DC decommission, disaster recovery

Correct Answer: D

AWS Snowball is a service that provides secure, rugged devices, so you can bring AWS computing storage capabilities to your edge environments, and transfer data into and out of AWS. Those rugged devices are commonly referred to as AWS Snowball or AWS Snowball Edge devices.

74) Which AWS shared responsibility controls are shared? (Select two.)

- •A. Awareness and training
- B. Patching of Amazon RDS
- •C. Configuration management
- D. Physical and environmental controls
- E. Service and communications protection or security

75) After a single Availability Zone service disruption, a corporation must guarantee that the endpoint for a database instance stays the same. The program must continue database operations without human intervention from an administrator.

Correct Answer: AC

How are these stipulations to be met?

- A. Use multiple Amazon Route 53 routes to the standby database instance endpoint hosted on AWS Storage Gateway.
- •B. Configure Amazon RDS Multi-Availability Zone deployments with automatic failover to the standby.
- C. Add multiple Application Load Balancers and deploy the database instance with AWS Elastic Beanstalk.
- D. Deploy a single Network Load Balancer to distribute incoming traffic across multiple Amazon CloudFront origins.

Correct Answer: B

76) Which costs must be addressed when comparing AWS Cloud vs on-premises Total Cost of Ownership? (Select two.)

- A. Software development
- B. Project management
- C. Storage hardware
- D. Physical servers
- E. Antivirus software license

Correct Answer: CD

77) Which costs must be addressed when comparing AWS Cloud vs on-premises Total Cost of Ownership? (Select two.)

- A. Physical storage hardware
- B. Operating system administration
- C. Network infrastructure of data center
- D. Project management
- E. Database schema development

Correct Answer: AC

78) Which AWS products anticipate future AWS expenses automatically?

- A. AWS Support Center
- B. AWS Total Cost of Ownership (TCO) Calculator
- C. AWS Simple Monthly Calculator
- ▶ •D. Cost Explorer

Correct Answer: D

AWS Cost Explorer has an easy-to-use interface that lets you visualize, understand, and manage your AWS costs and usage over time.

79) Which functionality may be utilized to prevent inadvertent overwrites or deletions of Amazon S3 buckets?

- A. Lifecycle policy
- B. Object versioning
- C. Server-side encryption
- D. Bucket ACL



Correct Answer: B

S3 Object Versioning is one of the most salient features in Amazon S3. Object versioning is used to avoid unintended overwrite and deletion. Versioning is not enabled by default and this feature is used to keep multiple versions of objects at the same time in the bucket.

If we enable versioning on the bucket, we cannot delete an object directly. All versions remain in the bucket and a delete marker is introduced which becomes the current version. So if you need to delete an object, you need to remove that delete marker also. Existing objects in your bucket do not change and only future requests behavior changes. If we put an object retrieval request, the current version of the object will always return.

80) Amazon DynamoDB is used by a business in its AWS Cloud architecture.

Which of the following is a duty of the organization, according to the AWS shared responsibility model? (Select two.)

- A. Operating system patching and upgrades
- •B. Application of appropriate permissions with IAM tools
- •C. Configuration of data encryption options
- D. Creation of DynamoDB endpoints
- E. Infrastructure provisioning and maintenance

Correct Answer: BC

81) Which of the following areas is the customer's duty under the shared responsibility model? (Select two.)

- A. Firmware upgrades of network infrastructure
- •B. Patching of operating systems
- C. Patching of the underlying hypervisor
- D. Physical security of data centers
- •E. Configuration of the security group

Correct Answer: BE

82) Which AWS service or functionality is utilized by distributed applications to send text and email messages?

- 
- A. Amazon Simple Notification Service (Amazon SNS)
 - B. Amazon Simple Email Service (Amazon SES)
 - C. Amazon CloudWatch alerts
 - D. Amazon Simple Queue Service (Amazon SQS)

Correct Answer: A

Amazon Simple Notification Service (Amazon SNS) is a fully managed messaging service for both application-to-application (A2A) and application-to-person (A2P) communication.

The A2A pub/sub functionality provides topics for high-throughput, push-based, many-to-many messaging between distributed systems, microservices, and event-driven serverless applications. Using Amazon SNS topics, your publisher systems can fanout messages to a large number of subscriber systems, including Amazon SQS queues, AWS Lambda functions, HTTPS endpoints, and Amazon Kinesis Data Firehose, for parallel processing. The A2P functionality enables you to send messages to users at scale via SMS, mobile push, and email.

83) To achieve high availability, how many Availability Zones should computing resources be provided across?



- A. A minimum of one
- B. A minimum of two
- C. A minimum of three
- D. A minimum of four or more

84) Which of the following is AWS's obligation under the AWS shared responsibility model?

Correct Answer: B

- A. Data encryption in transit
- B. Firmware updates on hardware
- C. Operating system patching on Amazon EC2 instances
- D. Data encryption at rest

85) To boost availability, a user intends to create two more Amazon EC2 instances.

What should the user do?

- •A. Launch the instances across multiple Availability Zones in a single AWS Region.
- B. Launch the instances as EC2 Reserved Instances in the same AWS Region and the same Availability Zone.
- C. Launch the instances in multiple AWS Regions, but in the same Availability Zone.
- D. Launch the instances as EC2 Spot Instances in the same AWS Region, but in different Availability Zones.

86) What are the customer's duties under the AWS shared responsibility model?
(Select two.)

Correct Answer: A

- A. Physical and environmental security
- B. Physical network devices including firewalls
- C. Storage device decommissioning
- •D. Security of data in transit
- •E. Data integrity authentication

Correct Answer: DE

87) How can customers minimize the amount of time they spend patching their operating systems by migrating to the AWS Cloud? (Select two.)

- •A. Users can take advantage of managed services on AWS.
- B. Users can outsource operating system patching to the AWS Support team.
- C. AWS Professional Services will upgrade instances to the latest operating system versions.
- D. Users have the ability to use license-included Amazon EC2 instances.
- •E. Users can take advantage of AWS Systems Manager features.

Correct Answer: AE

88) In the case of an environmental disruption, a company needs to make sure its infrastructure is structured for fault tolerance and business continuity. Which parts of the AWS architecture should the organization replicate?

- A. Edge locations
- B. Availability Zones
- C. Regions
- D. Amazon Route 53



Correct Answer: C

89) According to which AWS cloud design guideline, systems should minimize their interdependence?

- A. Scalability
- B. Services, not servers
- C. Removing single points of failure
- D. Loose coupling



Correct Answer: D

90) Which tasks need the root user credentials for an AWS account? (Select two.)

- A. Creating an Amazon EC2 key pair
- B. Removing an IAM user from the administrators group
- C. Changing the AWS Support plan
- D. Creating an Amazon CloudFront key pair
- E. Granting an IAM user full administrative access

91) An administrator must fast install and begin utilizing a popular IT product.

What resources are available to the administrator?

- A. AWS Well-Architected Framework documentation
- B. Amazon CloudFront
- C. AWS CodeCommit
- D. AWS Quick Start reference deployments

Correct Answer: CE

Correct Answer: D

Quick Starts are built by AWS solutions architects and partners to help you deploy popular technologies on AWS, based on AWS best practices for security and high availability. These accelerators reduce hundreds of manual procedures into just a few steps, so you can build your production environment quickly and start using it immediately.

92) Which solution enables users in various AWS Regions to have the FASTEST application response times for frequently requested data?

- A. AWS CloudTrail across multiple Availability Zones
- B. Amazon CloudFront to edge locations
- C. AWS CloudFormation in multiple regions
- D. A virtual private gateway over AWS Direct Connect



Correct Answer: B

You can deliver content and decrease end-user latency of your web application using Amazon CloudFront. CloudFront speeds up content delivery by leveraging its global network of data centers, known as edge locations, to reduce delivery time by caching your content close to your end users. CloudFront fetches your content from an origin, such as an Amazon S3 bucket, an Amazon EC2 instance, an Amazon Elastic Load Balancing load balancer or your own web server, when it's not already in an edge location. CloudFront can be used to deliver your entire website or application, including dynamic, static, streaming, and interactive content.

93) Which qualities make AWS Cloud computing advantageous? (Select two.)

- A. A 100% service level agreement (SLA) for all AWS services
- B. Compute capacity that is adjusted on demand
- C. Availability of AWS Support for code development
- D. Enhanced security
- E. Increases in cost and complexity

Correct Answer: BD

94) A corporation anticipates a brief increase in internet traffic for their application. The program cannot be interrupted during the traffic spike. In addition, the organization must reduce costs while increasing flexibility. To achieve these needs, which Amazon EC2 instance type should the organization use?

- A. On-Demand Instances

95) An Amazon RDS database instance is deployed across several Availability Zones.

Which pillar of the AWS Well-Architected Framework is included in this strategy?

- A. Performance efficiency
 - B. Reliability
 - C. Cost optimization
 - D. Security
- 

Correct Answer: B

96) Amazon EC2, an Elastic Load Balancer, and Amazon RDS are all components of an architectural design.

What is the BEST method for estimating the monthly cost of this architecture?

- A. Open an AWS Support case, provide the architecture proposal, and ask for a monthly cost estimation.
- B. Collect the published prices of the AWS services and calculate the monthly estimate.
- C. Use the AWS Simple Monthly Calculator to estimate the monthly cost.
- D. Use the AWS Total Cost of Ownership (TCO) Calculator to estimate the monthly cost.



Correct Answer: C

97) According to the AWS shared responsibility model, which job is the customer's duty?

- A. Maintain the security of the AWS Cloud.
- B. Configure firewalls and networks.
- C. Patch the operating system of Amazon RDS instances.
- D. Implement physical and environmental controls.

98) On Amazon EC2, a business hosts a web application in a Docker container.

Correct Answer: B

Which of the following duties is AWS in charge of?

- A. Scaling the web application and services developed with Docker
- B. Provisioning or scheduling containers to run on clusters and maintain their availability
- C. Performing hardware maintenance in the AWS facilities that run the AWS Cloud
- D. Managing the guest operating system, including updates and security patches

Correct Answer: C

98) On Amazon EC2, a business hosts a web application in a Docker container.

Which of the following duties is AWS in charge of?

- A. Scaling the web application and services developed with Docker
- B. Provisioning or scheduling containers to run on clusters and maintain their availability
- C. Performing hardware maintenance in the AWS facilities that run the AWS Cloud
- D. Managing the guest operating system, including updates and security patches

99) What are the AWS Cloud's advantages? (Select two.)

- A. Fixed rate monthly cost
- B. No need to guess capacity requirements
- C. Increased speed to market
- D. Increased upfront capital expenditure
- E. Physical access to cloud data centers

Correct Answer: C

Correct Answer: BC

100) An Elastic Load Balancer, numerous Amazon EC2 instances, and Amazon RDS are used to run a web application on AWS.

Which security measures are AWS's responsibility? (Select two.)

- A. Running a virus scan on EC2 instances
- •B. Protecting against IP spoofing and packet sniffing
- •C. Installing the latest security patches on the RDS instance
- D. Encrypting communication between the EC2 instances and the Elastic Load Balancer
- E. Configuring a security group and a network access control list (NACL) for EC2 instances

Correct Answer: BC

<https://www.youtube.com/@atozaboutdata/>